

2. Pods

27 February 2022 14:42

What Are Pods?

- Basic building block in Kubernetes
 - One or more containers in a Pod
 - Pod containers all share a container network
 - One IP address per Pod
- k8s will take care of it.

What's in a Pod Declaration



Container image



Container ports



Container restart policy



Resource limits

when to restart
if fails.

CPU / memory

Manifest Files

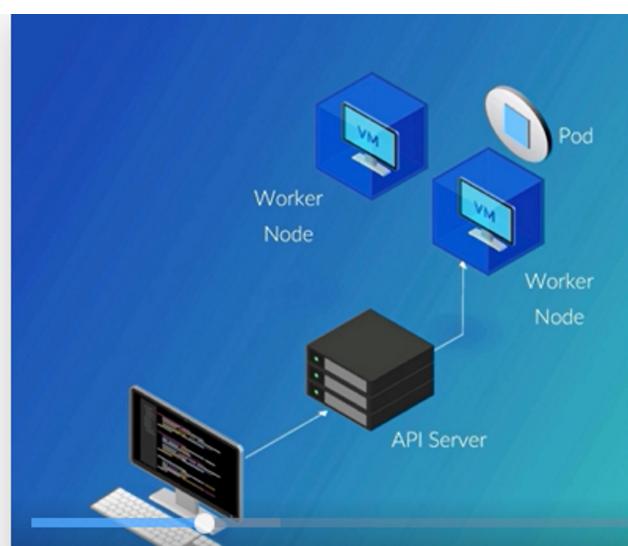
- Declare desired properties
- Manifests can describe all kinds of resource
- The spec contains resource-specific properties

1.1-basic_pod.yaml

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: mypod
5 spec:
6   containers:
7     - name: mycontainer
8       image: nginx:latest
```

Example of a basic manifest file

configuration
specific to a
resource.



Manifests in Action

- kubectl create sends manifest to Kubernetes API Server
- API Server does the following for Pod manifests
 - Select a node with sufficient resources
 - Schedule Pod onto node
 - Node pulls Pod's container image
 - Starts Pod's container

Why Use Manifests?



Can check into source control



Easy to share



Easier to work with

```
admin@Jamess-MBP Cloud Academy % minikube start
minikube v1.12.0 on Darwin 10.15.4
Using the hyperkit driver based on existing profile
Starting control plane node minikube in cluster minikube
Restarting existing hyperkit VM for "minikube" ...
Preparing Kubernetes v1.18.3 on Docker 19.03.12 ...
Verifying Kubernetes components...
Enabled addons: default-storageclass, storage-provisioner
Done! kubectl is now configured to use "minikube"
admin@Jamess-MBP Cloud Academy % kubectl get pods
No resources found in default namespace.
admin@Jamess-MBP Cloud Academy % cd src
admin@Jamess-MBP src % cat 1.1-basic_pod.yaml
apiVersion: v1
kind: Pod
metadata:
  name: mypod
spec:
  containers:
    - name: mycontainer
      image: nginx:latest
admin@Jamess-MBP src %
```

single container with image & name

```
admin@Jamess-MBP src % kubectl create -f 1.1-basic_pod.yaml
pod/mypod created
admin@Jamess-MBP src %
```

describe → to describe details.

Just like docker we need to tell the k8's which port we can be accessible.

```
Namespace: default
Priority: 0
Node: minikube/192.168.64.2
Start Time: Fri, 17 Jul 2020 12:09:23 -0400
Labels: <none>
Annotations: <none>
Status: Running
IP: 172.17.0.3
IPs:
  IP: 172.17.0.3
Containers:
  mycontainer:
    Container ID: docker://dd4aa6999d64bd425d97b09dbca5b93b72ec9c4314c1bad3f593f6579c707f
    Image: nginx:latest
    Image ID: docker-pullable://nginx@sha256:a9308a8b6974c967aebe868a186e5c205f4dc5423a5659f2f9599074bbcd
    Port: <none>
    Host Port: <none>
    State: Running
      Started: Fri, 17 Jul 2020 12:11:48 -0400
    Ready: True
    Restart Count: 0
    Environment: <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from default-token-7gjvg (ro)
Conditions:
:
```

```

admin@Jamess-MBP src % cat 1.2-port_pod.yaml
apiVersion: v1
kind: Pod
metadata:
  name: mypod
spec:
  containers:
    - name: mycontainer
      image: nginx:latest
      ports:
        - containerPort: 80
admin@Jamess-MBP src %

```

→ Port

K8's cannot update Ports for the running pod

C) we need to delete, update & re-create.

Kubectl delete pod mypod

(or)

Kubectl delete -f filename

```

Priority:      0
Node:          minikube/192.168.64.2
Start Time:    Fri, 17 Jul 2020 12:19:34 -0400
Labels:         <none>
Annotations:   <none>
Status:        Running
IP:            172.17.0.3
IPs:
  IP: 172.17.0.3
Containers:
  mycontainer:
    Container ID: docker://78edf6a17d358ff9cb33e501c74086a9cf56a80170c66b4a0e015c05
    Image:          nginx:latest
    Image ID:       docker-pullable://nginx@sha256:a93c8a0b0974c967aebe868a186e5cb5423a56559f2f9599074bbcd
    Port:           80/TCP
    Host Port:     0/TCP
    State:         Running
    Started:      Fri, 17 Jul 2020 12:19:36 -0400
    Ready:          True
    Restart Count: 0
    Environment:   <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from default-token-7gjvg (ro)
Conditions:
admin@Jamess-MBP src % curl 172.17.0.3:80

```

→ will not work

Because pod's ip is in the container network

```

Name:        mypod
Namespace:   default
Priority:    0
Node:        minikube/192.168.64.2
Start Time:  Fri, 17 Jul 2020 12:19:34 -0400
Labels:      <none> → key value pair like region.
Annotations: <none>
Status:      Running
IP:          172.17.0.3
IPs:

```

↳ helps to get pods in that region using label.

```

apiVersion: v1
kind: Pod
metadata:
  name: mypod
  labels:
    app: webserver
spec:
  containers:
    - name: mycontainer
      image: nginx:latest
      ports:
        - containerPort: 80
admin@Jamess-MBP src %

```

→ we can have multiple labels.

It's a good idea to set resource requests so that

It's a good idea to set resource requests so that API server knows to which node to schedule based on resource availability.

```
apiVersion: v1
kind: Pod
metadata:
  name: mypod
  labels:
    app: webserver
spec:
  containers:
  - name: mycontainer
    image: nginx:latest
    resources:
      requests:
        memory: "128Mi" # 128Mi = 128 mebibytes
        cpu: "500m"      # 500m = 500 milliCPUs (1/2 CPU)
      limits:
        memory: "128Mi"
        cpu: "500m"
    ports:
    - containerPort: 80
admin@James-MBP ~ %
```

```
Mounts:
  /var/run/secrets/kubernetes.io/serviceaccount from default-token-7gjvg (ro)
Conditions:
  Type          Status
  Initialized   True
  Ready         True
  ContainersReady True
  PodScheduled  True
Volumes:
  default-token-7gjvg:
    Type:       Secret (a volume populated by a Secret)
    SecretName: default-token-7gjvg
    Optional:   false
    QoS Class:  Guaranteed
    Node-Selectors: <none>
Tolerations:  node.kubernetes.io/not-ready:NoExecute for 300s
              node.kubernetes.io/unreachable:NoExecute for 300s
Events:
```

maximum req resource to schedule pod on node

maximum resource a node can assign to pod.

Pod will be guaranteed by service.