

7. Autoscaling

27 February 2022 15:29

Autoscaling :-

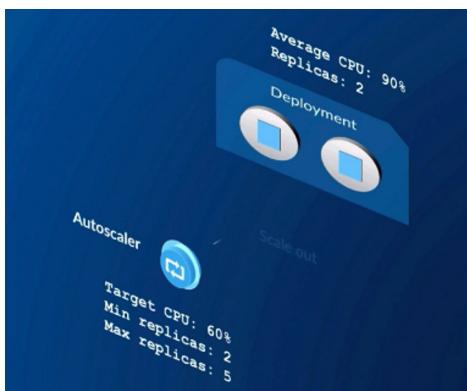
Autoscaling Deployments

 Scale automatically based on CPU utilization (or custom metrics)

 Set target CPU along with min and max replicas

 Target CPU is expressed as a percentage of the Pod's CPU request

If no CPU request is set autoscaling will not take any action.



Creates more replicas if average usage exceeds target.



→ Don't create more replicas than max set



→ Don't decrease less than min set replicas.

Autoscaling depends on metrics being collected on the cluster.

Metrics



Autoscaling depends on metrics being collected



Metrics Server is one solution for collecting metrics



Several manifest files are used to deploy Metrics Server
(<https://github.com/kubernetes-sigs/metrics-server>)

```
admin@Jamess-MBP src % kubectl apply -f metrics-server/
clusterrole.rbac.authorization.k8s.io/system:aggregated-metrics-reader created
clusterrolebinding.rbac.authorization.k8s.io/metrics-server:system:auth-delegator created
rolebinding.rbac.authorization.k8s.io/metrics-server-auth-reader created
apiservice.apiregistration.k8s.io/v1beta1.metrics.k8s.io created
serviceaccount/metrics-server created
deployment.apps/metrics-server created
service/metrics-server created
clusterrole.rbac.authorization.k8s.io/system:metrics-server created
clusterrolebinding.rbac.authorization.k8s.io/system:metrics-server created
admin@Jamess-MBP src %
```

```
admin@Jamess-MBP src % kubectl top pods -n deployments
NAME                               CPU(cores)   MEMORY(bytes)
app-tier-74cd4c68c9-h74qz        2m          45Mi
app-tier-74cd4c68c9-hcg26        1m          45Mi
app-tier-74cd4c68c9-p24sr        2m          45Mi
app-tier-74cd4c68c9-pscr         2m          45Mi
app-tier-74cd4c68c9-sfxgm        2m          46Mi
data-tier-8646dd765b-7f457       2m          2Mi
support-tier-997bc57fb-2dtq2     10m         2Mi
support-tier-997bc57fb-dlcv      11m         2Mi
support-tier-997bc57fb-gtd2d     12m         2Mi
support-tier-997bc57fb-kvvtr     11m         2Mi
support-tier-997bc57fb-mdlrp     12m         2Mi
admin@Jamess-MBP src %
admin@Jamess-MBP src %
```

Previous

```
-->
apiVersion: apps/v1
kind: Deployment
metadata:
  name: app-tier
  labels:
    app: microservices
    tier: app
spec:
  replicas: 1
  selector:
    matchLabels:
      tier: app
  template:
    metadata:
      labels:
        app: microservices
        tier: app
    spec:
      containers:
        - name: server
          image: lrakai/microservices:server-v1
          ports:
            - containerPort: 8080
      env:
        - name: REDIS_URL
          # Environment variable service discovery
          # Naming pattern:
          # IP address: <all_caps_service_name>_SERVICE_H
          # Port: <all_caps_service_name>_SERVICE_PORT
          # Named Port: <all_caps_service_name>_SERVICE_P
          value: redis://${DATA_TIER_SERVICE_HOST}:${DATA_T}
          # In multi-container example value was
          # value: redis://localhost:6379
admin@Jamess-MBP src %
```

updated

```
-->
apiVersion: apps/v1
kind: Deployment
metadata:
  name: app-tier
  labels:
    app: microservices
    tier: app
spec:
  replicas: 5
  selector:
    matchLabels:
      tier: app
  template:
    metadata:
      labels:
        app: microservices
        tier: app
    spec:
      containers:
        - name: server
          image: lrakai/microservices:server-v1
          ports:
            - containerPort: 8080
      resources:
        requests:
          cpu: 20m # 20 milliCPU / 0.02 CPU
      env:
        - name: REDIS_URL
          # Environment variable service discovery
          # Naming pattern:
          # IP address: <all_caps_service_name>_SERVICE_H
          # Port: <all_caps_service_name>_SERVICE_PORT
          # Named Port: <all_caps_service_name>_SERVICE_P
          value: redis://${DATA_TIER_SERVICE_HOST}:${DATA_T}
          # In multi-container example value was
          # value: redis://localhost:6379
admin@Jamess-MBP src %
```

→ It will deploy pods on to each node having min of 20 milliCPU

```
admin@Jamess-MBP src % kubectl create -f 6.1-app-tier-cpu_request.yaml -n deployments
Error from server (AlreadyExists): error when creating "6.1-app-tier-cpu_request.yaml": services "app-tier" already exists
Error from server (AlreadyExists): error when creating "6.1-app-tier_ccpu_request.yaml": deployments.apps "app-tier" already exists
admin@Jamess-MBP src %
admin@Jamess-MBP src %
admin@Jamess-MBP src %
```

→ We can delete & recreate

(or)

→ use `apply` command which updates the deployment.

... & create

(Q6)
 → use apply command which update the deployment.
 → use apply command which update the deployment.
 ✓ CKAD administrator course will have idea on diff between apply & create

```
admin@Jamess-MBP ~ % kubectl get -n deployments deployments app-tier
NAME      READY   UP-TO-DATE   AVAILABLE   AGE
app-tier  5/5     5           5           13m
```

```
admin@Jamess-MBP ~ % cat 6.2-autoscale.yaml
apiVersion: autoscaling/v1
kind: HorizontalPodAutoscaler
metadata:
  name: app-tier
  labels:
    app: microservices
    tier: app
spec:
  maxReplicas: 5
  minReplicas: 1
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: app-tier
  targetCPUUtilizationPercentage: 70
```

→ kind for autoscaling
 → lower & upper limit
 → 70%
 # Equivalent to
 # kubectl autoscale deployment app-tier --max=5 --min=1 --cpu-percent=70

version
 command for above manifest—

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
app-tier	1/1	1	1	19m
data-tier	1/1	1	1	19m
Support-tier	5/5	5	5	19m

→ decreased

admin@Jamess-MBP ~ % kubectl api-resources

→ gives all resources with their short names.

customresourcedefinitions	crd,crds	apiextensions.k8s.io	false	CustomResourceDefinition
apiservices		apiregistration.k8s.io	false	APIService
controllerrevisions		apps	true	ControllerRevision
deployments	ds	apps	true	Deployment
replicsets	rs	apps	true	ReplicaSet
statefulsets	sts	apps	true	StatefulSet
tokenreviews		authentication.k8s.io	false	TokenReview
localsubjectaccesreviews		authorization.k8s.io	true	LocalSubjectAccessReview
selfsubjectaccesreviews		authorization.k8s.io	false	SelfSubjectAccessReview
selfsubjectrulesreviews		authorization.k8s.io	false	SelfSubjectRulesReview
subjectaccesreviews		authorization.k8s.io	false	SubjectAccessReview
horizontalpodautoscalers	hpa	autoscaling	true	HorizontalPodAutoscaler
cronjobs		batch	true	CronJob
jobs		batch	true	Job
certificatesigningrequests	csr	certificates.k8s.io	false	CertificateSigningRequest
leases		coordination.k8s.io	true	Lease
endpointslices		discovery.k8s.io	true	EndpointSlice
events	ev	events.k8s.io	true	Event
leases		lease.k8s.io	true	Lease

→ short name

```
admin@Jamess-MBP ~ % kubectl describe -n deployments hpa
Name:          app-tier
Namespace:    default
Labels:        app=app-tier
               app=microservices
               tier=app
Annotations: 
CreationTimestamp:  Tue, 04 Aug 2020 12:10:16 +1200
Reference:    Deployment/app-tier
Metrics:      resource cpu on pods  (as a percentage of request):  40% (8m) / 70%
Min replicas: 1
Max replicas: 5
Deployment pods:       1 current / 1 desired
Conditions:
  Type      Status  Reason
  ----      ----  -----
  AbleToScale  True    ReadyForNewScale  recommended size matches current size
  ScalingActive  True    ValidMetricFound  the HPA was able to successfully calculate a replica count from cpu resource utilization (percent)
  age of request
  ScalingLimited  False   DesiredWithinRange  the desired count is within the acceptable range
Events:
  Type      Reason
  ----      -----
  Normal   SuccessfulRescale  107s  horizontal-pod-autoscaler  New size: 2; reason: All metrics below target
  Normal   SuccessfulRescale  77s   horizontal-pod-autoscaler  New size: 1; reason: All metrics below target
```

```
admin@Jamess-MBP src % kubectl get -n deployments hpa
NAME      REFERENCE   TARGETS   MINPODS   MAXPODS   REPLICAS   AGE
app-tier  Deployment/app-tier  40%/70%    1          5          1          7m47s
admin@Jamess-MBP src %
```

```
admin@Jamess-MBP src %
admin@Jamess-MBP src % kubectl edit -n deployments hpa
```

(we can update manifest file & use apply)
 (or) (use edit command):
 open in VT & we can edit the config like replicas if we want.

```
f:apiVersion: v1
f:kind: HorizontalPodAutoscaler
f:targetCPUUtilizationPercentage: 70
manager: kubectl
operation: Update
time: "2020-08-04T00:10:16Z"
- apiVersion: autoscaling/v1
  fieldType: FieldsV1
  fields:
    f:metadata:
      f:annotations:
        :: {}
      f:autoscaling.alpha.kubernetes.io/conditions: {}
      f:autoscaling.alpha.kubernetes.io/current-metrics: {}
    f:status:
      f:currentCPUUtilizationPercentage: 0
      f:currentReplicas: 0
      f:desiredReplicas: 0
      f:lastScaleTime: {}
  manager: kubectl-controller-manager
  operation: Update
  time: "2020-08-04T00:18:07Z"
  name: app-tier
  namespace: deployments
  resourceVersion: "4191"
  selfLink: "/apis/autoscaling/v1/namespaces/deployments/horizontalpodautoscalers/app-tier"
  uid: d81180be-f4aa-432d-b5cc-67de68535e62
spec:
  maxReplicas: 5
  minReplicas: 2
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: app-tier
  targetCPUUtilizationPercentage: 70
status:
```

In VT

```
admin@Jamess-MBP src %
admin@Jamess-MBP src % kubectl edit -n deployments hpa
horizontalpodautoscaler.autoscaling/app-tier edited
```

```
Every 1.0s: kubectl get -n deployments deployment
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
app-tier	2/2	2	2	24m
data-tier	1/1	1	1	24m
support-tier	5/5	5	5	24m

CONCLUSION

- Autoscaling depends on metrics being collected
- Pods must have CPU requests
- horizontalPodAutoscaler (hpa) is configured with target CPU, and min and max replicas
- kubectl apply to update resources
- kubectl edit to edit and apply

