

14. Multiple Container Patterns

01 March 2022 01:18

Multi-Container Patterns



extra level of abstraction for containers.

Why Pods?

→ containers are not enough for Kubernetes to effectively manage



Kubernetes needs additional information



Simplifies using different underlying container runtimes

(e.g. - docker / rocket)



Co-locate tightly coupled containers without packaging them as a single image

CloudAcademy



use a helper container to assist the

Sidecar Pattern Primary container



Uses a helper container to assist a primary container



Commonly used for logging, file syncing, watchers



Benefits include leaner main container, failure isolation, independent update cycles

all these can be accomplished by using a sidecar than using the main container to do these things.

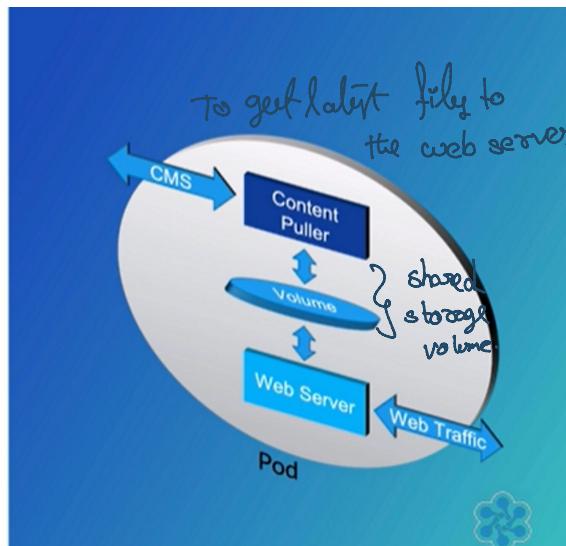
Main container can do the traffic.

CloudAcademy



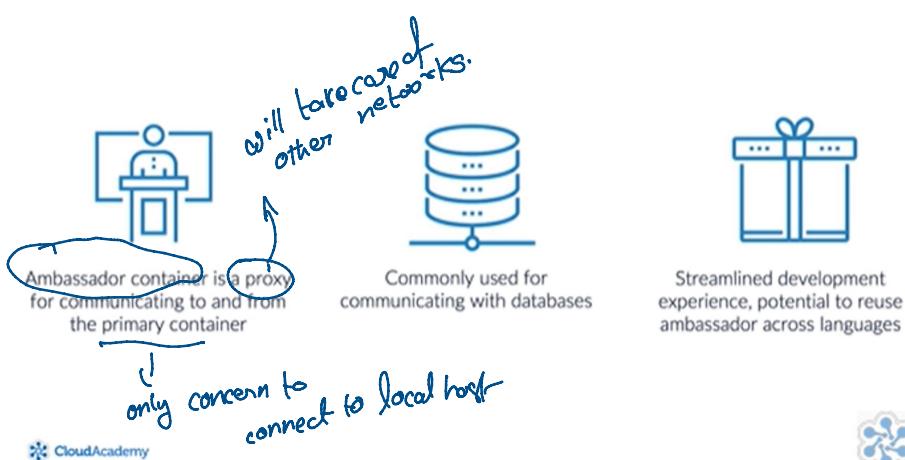
File Sync Sidecar Example

- Primary container: Web Server
- Sidecar: Content Puller
- Content Puller syncs with content management system (CMS)
- Web Server serves content
- The content is synced using a volume



CloudAcademy

Ambassador Pattern



CloudAcademy

Database Ambassador Example

- Primary container: Web App
- Ambassador: Database Proxy
- Web App handles requests
- Database requests are sent to the Database Proxy over localhost
- Database Proxy then forwards the requests to the appropriate database
- Possibly sharding the requests



CloudAcademy

Adapter Pattern

Adapter Pattern

simple for outside
coastal.



Adapters present a standardized interface across multiple Pods



{ Commonly used for normalizing
output logs and monitoring data }
across the pods



Adapts third-party software to
meet your needs

In this demonstration, we'll see how to implement the adapter pattern.

```
ubuntu@ip-10-0-128-5:~$ kubectl exec -n legacy app -- cat /metrics/raw.txt
Thu Mar 28 23:14:57 UTC 2019
Mem: 843844K used, 136912K free, 1180K shrd, 24920K buff, 410652K cached
CPU: 0% usr 0% sys 0% nic 100% idle 0% io 0% irq 0% sirq
Load average: 0.01 0.02 0.08 1/334 143
  PID  PPID USER      STAT  VSZ %VSZ CPU %CPU COMMAND
    1     0 root      S    1588   0%  0% /bin/sh -c mkdir /metrics; while true; do
date > /metrics/raw.txt; top -n 1 -b > /metrics/raw.txt; sleep 5; done
   43     1 root      R    1528   0%  0% top -n 1 -b
} we only CPU
```

↳ This is where we can use adaptive

```
ubuntu@ip-10-0-128-5:~$ vi no-adapter.yaml
```

```

1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: adapted-pod
5 spec:
6   containers:
7     - name: app
8       image: alpine:3.9.2
9       command: ["/bin/sh", "-c"]
10      args:
11        - while true; do
12          date > /metrics/raw.txt;
13          top -n 1 -b >> /metrics/raw.txt;
14          sleep 5;
15        done
16      # Share metric data between containers
17      volumeMounts:
18        - name: metrics
19          mountPath: /metrics
20      - name: adapter
21        image: httpd:2.4.38-alpine
22        command: ["/bin/sh", "-c"]
23        # Adapt the legacy output to standard monitoring format
24        args:
25          - while true; do
26            date=$(head -1 /metrics/raw.txt);
27
28 "adapter.yaml" 38L, 1044C

```

the legacy app remain same
with one diff of volume so that raw metrics can be shared with adaptive container.

} → adaptive container,



```

14     sleep 5;
15   done
16   # Share metric data between containers
17   volumeMounts:
18     - name: metrics
19       mountPath: /metrics
20   - name: adapter
21     image: httpd:2.4.38-alpine
22     command: ["/bin/sh", "-c"]
23     # Adapt the legacy output to standard monitoring format
24     args:
25       - while true; do
26         date=$(head -1 /metrics/raw.txt); ✓
27         memory=$(head -2 /metrics/raw.txt | tail -1 | grep -o -E '\d+\w*' | head -1)
28         cpu=$(head -3 /metrics/raw.txt | tail -1 | grep -o -E '\d+\%' | head -1);
29         echo "{\"date\":\"$date\", \"memory\":\"$memory\", \"cpu\":\"$cpu\"}" > /metrics/adapter.json;
30       sleep 5;
31     done
32   # Share metric data between containers
33   volumeMounts:
34     - name: metrics
35       mountPath: /metrics
36   volumes:
37     - name: metrics
38       emptyDir: ✓ means volume will have same life time as the pod.

```

} json output every 5 seconds



38,4 Bot

↳ there is no need of persistent volume.

```

ubuntu@ip-10-0-128-5:~$ kubectl create -f adapter.yaml
pod/adapter-pod created
ubuntu@ip-10-0-128-5:~$ 

```

A

```

ubuntu@ip-10-0-128-5:~$ kubectl exec adapted-pod --container=adapter -- cat /metrics/adapter.json
{"date":"Thu Mar 28 23:34:37 UTC 2019","memory":"839892K","cpu":"0%"} 
ubuntu@ip-10-0-128-5:~$ 

```

} json in the metrics

```

ubuntu@ip-10-0-128-5:~$ kubectl exec adapted-pod --container=adapter -- cat /metrics/adapter.json
{"date":"Thu Mar 28 23:34:42 UTC 2019","memory":"839620K","cpu":"0%"} 
ubuntu@ip-10-0-128-5:~$ 

```

} we can use
webserver to give the
metric as a rest API
or configuring the adapter to
push json to metric aggregator
system.

CONCLUSION

- Explained why Pods

From json to metric gathering system.

CONCLUSION

- Explained why Pods
- Sidecar, ambassador, and adapter patterns
- Demo implementation of adapter pattern