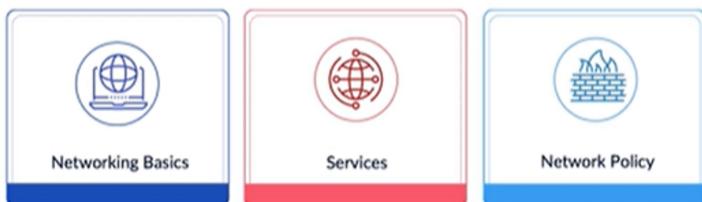


## 15. Networking

01 March 2022 01:29

### Networking

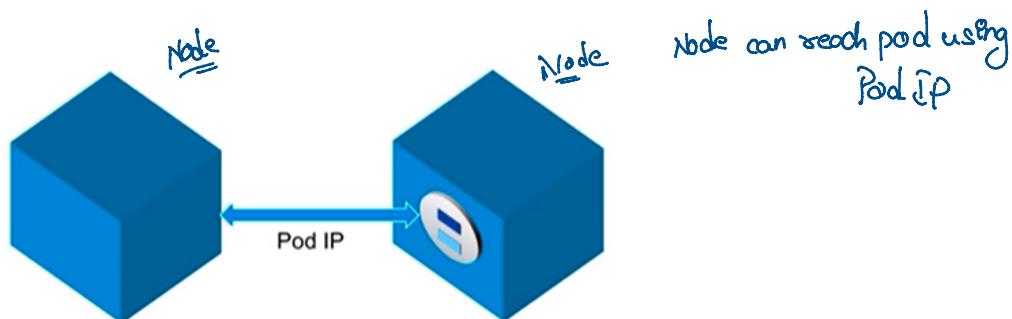


### Kubernetes Networking Basics

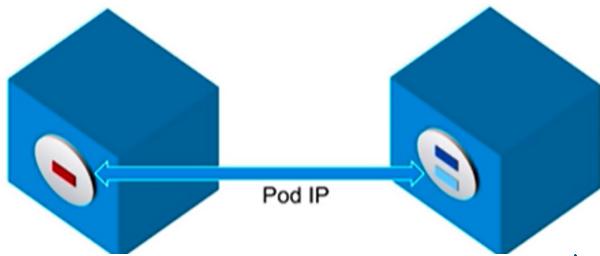
→ each pod one unique IP  
→ container can use local host



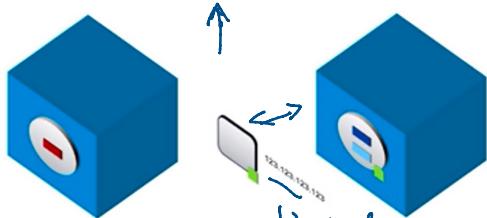
### Kubernetes Networking Basics



## Kubernetes Networking Basics



- o Pod can reach other pod using pod IP
- o killed and restarted on diff IP
- o replace diff IP



service can be used based on labels.

o service can use env variable  
(or) DNS

cluster IP only reachable within the cluster.

Node port → is opened in every node of the cluster any seq.  
to node port will go to the cluster IP.

Load balancer → exposed to outside by using cloud provider load balancer.

## Kubernetes Network Policy



Rules for controlling network access to Pods



Similar to security groups controlling access to virtual machines



Scoped to Namespace

↳ we can create network policy independently to each namespace.

## Network Policy Caveat



Kubernetes network plugin must support Network Policy

with no

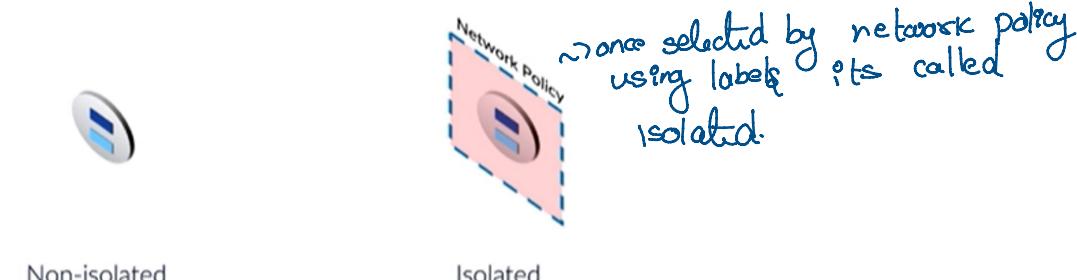


Kubernetes network plugin must support Network Policy

in per cluster

otherwise though we successfully create Network policy  
there is no way to enforce the policy <sup>↑ with no errors.</sup>

## Non-Isolated vs. Isolated Pods



```
ubuntu@ip-10-0-128-5:~$ kubectl get pods -n kube-system
NAME                                         READY   STATUS    RESTARTS   AGE
calico-etcd-zb28f                           1/1    Running   2          11d
calico-kube-controllers-5dcc459c88-74h7v   1/1    Running   3          11d
calico-node-cftgz                          2/2    Running   2          12m
calico-node-dbv2z                          2/2    Running   1          12m
calico-node-ms9qb                          2/2    Running   7          11d
coredns-86c58d9df4-cmc8n                  1/1    Running   2          11d
coredns-86c58d9df4-n2vpg                  1/1    Running   2          11d
etcd-ip-10-0-5-100.us-west-2.compute.internal 1/1    Running   2          11d
kube-apiserver-ip-10-0-5-100.us-west-2.compute.internal 1/1    Running   2          11d
kube-controller-manager-ip-10-0-5-100.us-west-2.compute.internal 1/1    Running   2          11d
kube-proxy-8tqkj                           1/1    Running   0          12m
kube-proxy-cd7rc                           1/1    Running   2          11d
kube-proxy-d9sk2                           1/1    Running   0          12m
kube-scheduler-ip-10-0-5-100.us-west-2.compute.internal 1/1    Running   2          11d
kubernetes-dashboard-59756dbcc8-whkmn      1/1    Running   2          11d
metrics-server-68d85f76bb-g4dhf           1/1    Running   0          10d
ubuntu@ip-10-0-128-5:~$
```

```
ubuntu@ip-10-0-128-5:~$ kubectl get pods -n network-policy -L app,region
NAME     READY   STATUS    RESTARTS   AGE   APP   REGION
client-1  1/1    Running   0          97s   client  us-east
client-2  1/1    Running   0          97s   client  us-west
server   1/1    Running   0          97s   server  us-central
ubuntu@ip-10-0-128-5:~$
```

} two clients in diff regions sends requests to server

```
ubuntu@ip-10-0-128-5:~$ kubectl logs -n network-policy --tail=1 -f client-1
Received request
Received request
Received request
Received request
^C
ubuntu@ip-10-0-128-5:~$ kubectl logs -n network-policy --tail=1 -f client-2
Received request
Received request
```

} server receiving the requests

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: allow-us-east
  namespace: network-policy
spec:
  # Which Pods the policy applies to
  podSelector:
    matchLabels:
      app: server
  policyTypes:
  - Ingress
  - Egress
  ingress:
  - from:
    # Supports podSelector, namespaceSelector, and ipBlock
    - podSelector:
        matchLabels:
          region: us-east
        ports: # If not present allows all ports
        - protocol: TCP
          port: 8888
  egress:
  - to: # Allows all traffic when empty
```

} network policies are available in the networking api's

} allowing us-east

} if pod selector empty then the policy applies to all the pods.

} if pods selector empty then the policy applies to all the pods.

} we can include one or both.

} supports

} allows to all outbound traffic.

} if from is omitted traffic from all sources is allowed

} if port is omitted traffic from all ports is allowed

```
ubuntu@ip-10-0-128-5:~$ kubectl logs -n network-policy --tail=0 -f client-1
Received request
Received request
Received request
^C
```

} only us east requests are coming.

```

ubuntu@ip-10-0-128-5:~$ kubectl logs -n network-policy --tail=0 -f client-1
Received request
Received request
Received request
^C
ubuntu@ip-10-0-128-5:~$ kubectl logs -n network-policy --tail=0 -f client-2
^C
ubuntu@ip-10-0-128-5:~$ 

```

} only requests are coming.

```

apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: block-one-ip
  namespace: network-policy
spec:
  podSelector:
    matchLabels:
      app: server
  policyTypes:
    - Egress
  egress:
    - to:
        - ipBlock:
            cidr: 0.0.0.0/0
    except:
      - 192.168.134.72/32

```

} set range of allowed IP's  
→ This means all IP addresses are allowed.

optional | Block list.

- This is IP of client 1 for demonstration purpose.
- In general we should use labels to refer to pods in your cluster; because IP address changes, label remains the same.

```

ubuntu@ip-10-0-128-5:~$ kubectl logs -n network-policy --tail=0 -f client-2
^C
ubuntu@ip-10-0-128-5:~$ kubectl logs -n network-policy --tail=0 -f client-1
Received request
Received request
Received request
Received request
^C
ubuntu@ip-10-0-128-5:~$ clear

```

} How we are getting response though we blocked.

If one policy allows the traffic though other deny it.  
The traffic is allowed.

1st policy allowed enough to all } allowed to client 1  
2nd denied to client 1

```

ubuntu@ip-10-0-128-5:~$ kubectl logs -n network-policy --tail=0 -f client-1
Received request
Received request
Received request
Received request
^C
ubuntu@ip-10-0-128-5:~$ 

```

if we delete the old networkpolicy & keep only one, still we see traffic coming in.

Because the rules apply to new connection. As old connection still exists traffic is still flowing.

```

ubuntu@ip-10-0-128-5:~$ kubectl exec -n network-policy server -it ping 192.168.134.72
PING 192.168.134.72 (192.168.134.72): 56 data bytes
^C
--- 192.168.134.72 ping statistics ---
6 packets transmitted, 0 packets received, 100% packet loss
command terminated with exit code 1
ubuntu@ip-10-0-128-5:~$ kubectl exec -n network-policy server -it ping 192.168.233.198
PING 192.168.233.198 (192.168.233.198): 56 data bytes
64 bytes from 192.168.233.198: seq=0 ttl=63 time=0.076 ms
64 bytes from 192.168.233.198: seq=1 ttl=63 time=0.058 ms
64 bytes from 192.168.233.198: seq=2 ttl=63 time=0.057 ms
64 bytes from 192.168.233.198: seq=3 ttl=63 time=0.058 ms
^C
--- 192.168.233.198 ping statistics ---
4 packets transmitted, 4 packets received, 0% packet loss
round-trip min/avg/max = 0.057/0.062/0.076 ms
ubuntu@ip-10-0-128-5:~$ 

```

} pinging from server to client 1  
we see there is no response because trying to establish a new connection here.

# CONCLUSION

- Networking review
- Types of Kubernetes services
- Network Policy