

Kubernetes

Wednesday, February 23, 2022 1:01 AM

cluster → more than one node / compute / system → come together in a network to perform task

In k8 we can have clusters from cloud / legacy (vpc)

- Controlled | Manged (version is depending on the provider)
we control but time consuming
AKS → azure
EKS → aws

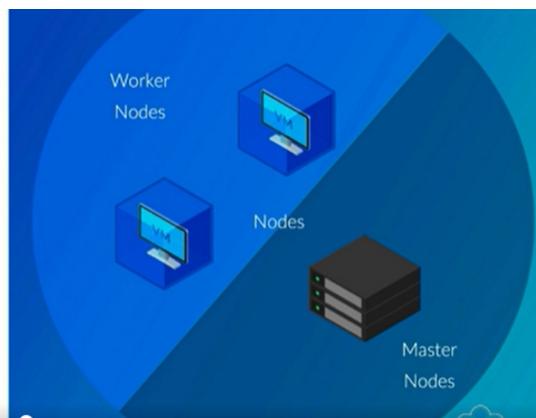
other options similar to k8 is docker swarm mode / EKS / DO / GKE

- GitHub - cloudacademy/intro-to-k8s: Assets used for the production of the Introduction to Kubernetes course on Cloud Academy

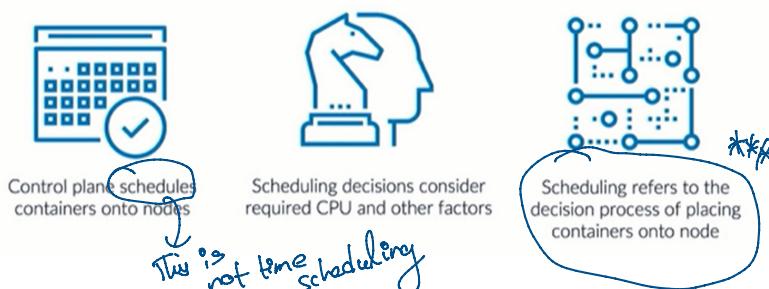
- GitHub - kubernetes-sigs/metrics-server: Scalable and efficient source of container resource metrics for Kubernetes built-in autoscaling pipelines.

Kubernetes Architecture

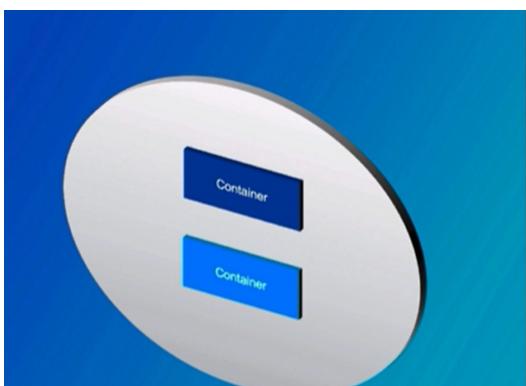
- Cluster refers to all of the machines collectively and can be thought of as the entire running system
- Nodes are the machines in the cluster
- Nodes are categorized as workers or masters
- Worker nodes include software to run containers managed by the Kubernetes control plane
- Master nodes run the control plane
- The control plane is a set of APIs and software that Kubernetes users interact with
- The APIs and software are referred



Scheduling



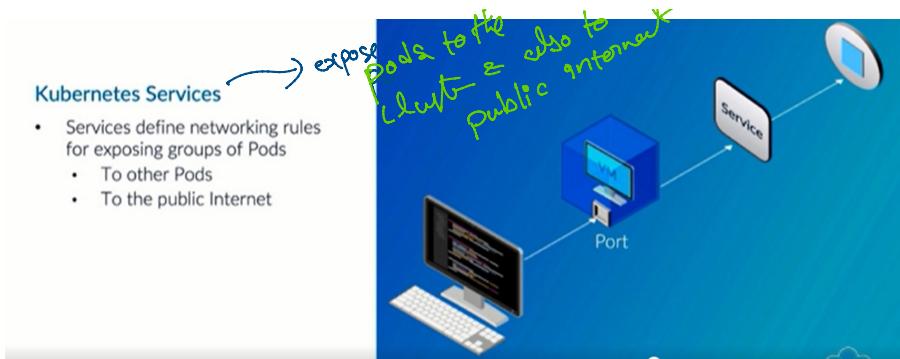
Pods → Group of containers (one or more)
↳ all containers in pod run on the same node.



Kubernetes Pods

- Groups of containers
- Pods are the smallest building block in Kubernetes

Kubernetes Services → expose pods to the outside & also to the internal



Kubernetes Deployments → control rollback & rollout of pods.

- Manage deploying configuration changes to running Pods
- Horizontal scaling

→ Kubernetes API server

 Modify cluster state information by sending requests to the Kubernetes API Server

 The API Server is a master component that acts as the frontend for the cluster

Interacting with Kubernetes: Method 1

REST API



It is possible but not common to work directly with the API server

You might need to if there is no Kubernetes client library for your programming language

Interacting with Kubernetes: Method 2

Client Libraries



Handles authenticating and managing individual REST API requests and responses



Kubernetes maintains official client libraries



Community-maintained libraries when no official library exists

Interacting with Kubernetes: Method 3

kubectl



Issue high-level commands that are translated into REST API calls



Works with local and remote clusters

* * * Most commonly used method is this one.

kubectl



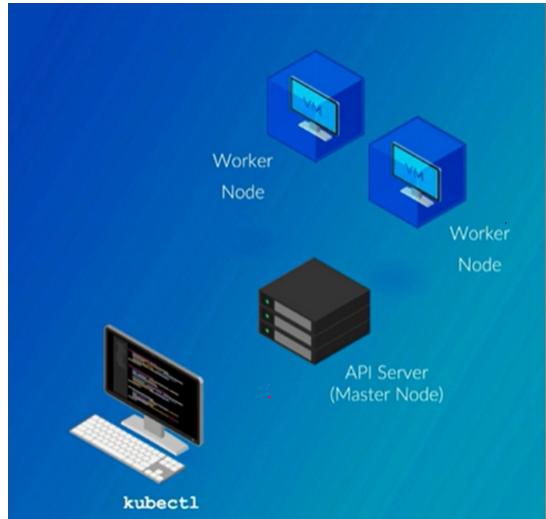
Kubernetes success correlates with kubectl skill



Manages all different types of Kubernetes resources, and provides debugging and introspection features



kubectl commands follow an easy to understand pattern



Example kubectl commands

- kubectl create to create resources (Pods, Services, etc.)
- kubectl delete to delete resources
- kubectl get to get list of resources of a given type
 - kubectl get pods
- kubectl describe to print detailed info about a resource(s)
 - kubectl describe pod server
- kubectl logs to print container logs

→ commands can be used directly (or)
you can also refer to a file called manifest
↑
where all resources names manifest
are declared

Interacting with Kubernetes: Method 4

Web Dashboard → optional
not all clients will have this

Name	Node	Status	Replicas	Age
mysql-4	ip-10-0-19-165.us-west-2.compute.internal	Running	0	7 minutes
mysql-3	ip-10-0-2-43.us-west-2.compute.internal	Running	0	7 minutes
mysql-2	ip-10-0-2-43.us-west-2.compute.internal	Running	0	9 minutes
mysql-1	ip-10-0-19-154.us-west-2.compute.internal	Running	0	20 minutes
mysql-5	ip-10-0-8-120.us-west-2.compute.internal	Running	0	21 minutes

What Are Pods?



Basic building block in Kubernetes

One or more containers in a Pod

Pod containers all share a container network

One IP address per Pod

→ pods will take care of it.

What's in a Pod Declaration



What's in a Pod Declaration

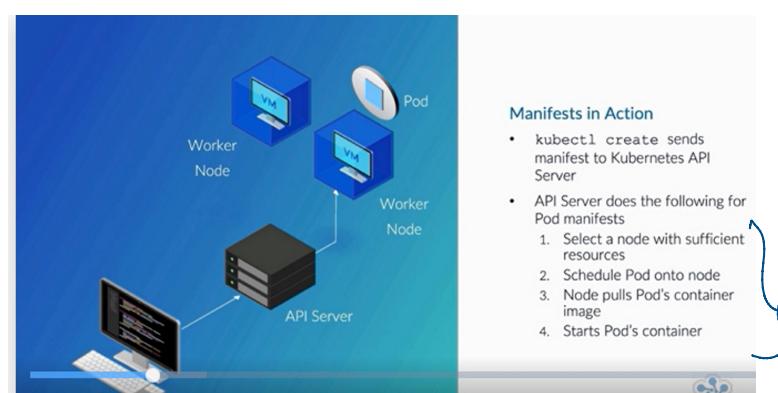
Manifest Files

- Declare desired properties
- Manifests can describe all kinds of resource
- The spec contains resource-specific properties

Example of a basic manifest file

```
1.1-basic_pod.yaml
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: mypod
5 spec:
6   containers:
7     - name: mycontainer
8       image: nginx:latest
```

configuration specific to a specific source.



Why Use Manifests?

Can check into source control

Easy to share

Easier to work with

```
admin@Jamess-MBP Cloud Academy % minikube start
minikube v1.12.0 on Darwin 10.15.4
Using the hyperkit driver based on existing profile
Starting control plane node minikube in cluster minikube
Restarting existing hyperkit VM for "minikube" ...
Preparing Kubernetes v1.18.3 on Docker 19.03.12 ...
Verifying Kubernetes components...
Enabled addons: default-storageclass, storage-provisioner
Done! kubectl is now configured to use "minikube"
admin@Jamess-MBP Cloud Academy % kubectl get pods
No resources found in default namespace.
admin@Jamess-MBP Cloud Academy % cd src
admin@Jamess-MBP src % cat 1.1-basic_pod.yaml
apiVersion: v1
kind: Pod
metadata:
  name: mypod
spec:
  containers:
    - name: mycontainer
      image: nginx:latest
admin@Jamess-MBP src %
```

single container with image & name

describe → to describe details.

create using a file

Just like docker we need to tell the k8's which port we can be accessible.

```

Namespace: default
Priority: 0
Node: minikube/192.168.64.2
Start Time: Fri, 17 Jul 2020 12:09:23 -0400
Labels: <none>
Annotations: <none>
Status: Running
IP: 172.17.0.3
IPs:
  IP: 172.17.0.3
Containers:
  mycontainer:
    Container ID: docker://dd4aa6999d64bd425d97b09dbca5b93b72ec9c4314c1bad3f593f6579c70f
      Image: nginx:latest
      Image ID: docker-pullable://nginx@sha256:a93c8a0b0974c967aebe868a186e5c205f4dc5423a5659f2f9599874bbcd
      Port: <none>
      Host Port: <none>
      State: Running
      Started: Fri, 17 Jul 2020 12:11:48 -0400
      Ready: True
      Restart Count: 0
      Environment: <none>
      Mounts:
        /var/run/secrets/kubernetes.io/serviceaccount from default-token-7gjvg (ro)
Conditions:
:|
```

```

admin@Jamess-MBP src % cat 1.2-port_pod.yaml
apiVersion: v1
kind: Pod
metadata:
  name: mypod
spec:
  containers:
    - name: mycontainer
      image: nginx:latest
      ports:
        - containerPort: 80 → Port
admin@Jamess-MBP src %
```

I⁸'s cannot update Ports for the running pod
 (c) we need to ~~edit~~, update & recreate.

Kubectl delete pod my pod

(o8)
 kubectl delete -f filename

```

Priority: 0
Node: minikube/192.168.64.2
Start Time: Fri, 17 Jul 2020 12:19:34 -0400
Labels: <none>
Annotations: <none>
Status: Running
IP: 172.17.0.3
IPs:
  IP: 172.17.0.3
Containers:
  mycontainer:
    Container ID: docker://78edf6a17d358ff9cb33e501c74086a9cf56a80170c66b4a0e0c5c65
      Image: nginx:latest
      Image ID: docker-pullable://nginx@sha256:a93c8a0b0974c967aebe868a186e5c5423a5659f2f9599874bbcd
      Port: 80/TCP
      Host Port: 80/TCP
      State: Running
      Started: Fri, 17 Jul 2020 12:19:36 -0400
      Ready: True
      Restart Count: 0
      Environment: <none>
      Mounts:
        /var/run/secrets/kubernetes.io/serviceaccount from default-token-7gjvg (ro)
Conditions:
admin@Jamess-MBP src % curl 172.17.0.3:80 → will not work.
```

Because pod's ip is in the container network.

```

Name: mypod
Namespace: default
Priority: 0
Node: minikube/192.168.64.2
Start Time: Fri, 17 Jul 2020 12:19:34 -0400
Labels: <none> → key value pair like regions.
Annotations: <none>
Status: Running
IP: 172.17.0.3
IPs:
```

↳ helps to get pods in that region using label.

```

apiVersion: v1
kind: Pod
metadata:
  name: mypod
  labels:
    app: webserver
spec:
  containers:
    - name: mycontainer
      image: nginx:latest
      ports:
        - containerPort: 80
admin@Jamess-MBP src %
```

→ we can have multiple labels.

It's a good idea to set resource requests so that API server knows to which node to schedule based on resource availability.

↳ node)

```

apiVersion: v1
kind: Pod
metadata:
  name: mypod
  labels:
    app: webserver
spec:
  containers:
  - name: mycontainer
    image: nginx:latest
  resources:
    requests:
      memory: "128Mi" # 128Mi = 128 mebibytes
      cpu: "500m" # 500m = 500 milliCPUs (1/2 CPU)
    limits:
      memory: "128Mi"
      cpu: "500m"
  ports:
  - containerPort: 80
admin@Jamess-MBP src %

```

minimum req resource to schedule pod on node

maximum resource a node can assign to pod.

Pod will be guaranteed this service.

```

Mounts:
  /var/run/secrets/kubernetes.io/serviceaccount from default-token-7gjvg (ro)
Conditions:
  Type          Status
  Initialized   True
  Ready         True
  ContainersReady True
  PodScheduled  True
Volumes:
  default-token-7gjvg:
    Type:       Secret (a volume populated by a Secret)
    SecretName: default-token-7gjvg
    Optional:   false
    QoS Class:  Guaranteed
    Node-Selectors: <none>
    Toleration: node.kubernetes.io/not-ready:NoExecute for 300s
    node.kubernetes.io/unreachable:NoExecute for 300s
Events:

```

To overcome networking issues k8's provides services

Services

- A service defines networking rules for accessing Pods in the cluster and from the internet
- Use labels to select a group of Pods
- Service has a fixed IP address
- Distribute requests across Pods in the group

```

apiVersion: v1
kind: Service
metadata:
  labels:
    app: webserver
    name: webserver
spec:
  ports:
  - port: 80
    selector:
      app: webserver
    type: NodePort

```

selector define the labels to match the pod against.

values of pod's container port

How to expose the service.

NodePort → allocates Port over this service on each node in

the cluster

By doing this you can reach this service from any node by sending request to this service.

```

admin@Jamess-MBP src % kubectl create -f 2.1-web_service.yaml
service/webserver created
admin@Jamess-MBP src % kubectl get services
NAME      TYPE        CLUSTER-IP   EXTERNAL-IP   PORT(S)   AGE
kubernetes  ClusterIP  10.96.0.1   <none>        443/TCP   6d18h
webserver  NodePort   10.100.37.149  <none>        80:32337/TCP 5s
admin@Jamess-MBP src %

```

private IP for each service.

not available for NodePort type

It's automatically allocated port

normally between 30,000 & 32,767

```

admin@Jamess-MBP src % kubectl get services
NAME      TYPE      CLUSTER-IP   EXTERNAL-IP   PORT(S)      AGE
kubernetes  ClusterIP  10.96.0.1    <none>        443/TCP     6d18h
webserver   NodePort   10.100.37.149  <none>        80:32337/TCP  5s
admin@Jamess-MBP src % kubectl describe service webserver
Name:           webserver
Namespace:      default
Labels:         app=webserver
Annotations:   <none>
Selector:       app=webserver
Type:          NodePort
IP:            10.100.37.149
Port:          <unset>  80/TCP
TargetPort:    80/TCP
NodePort:     <unset>  32337/TCP
Endpoints:    172.17.0.3:80
Session Affinity: None
External Traffic Policy: Cluster
Events:        If there were multiple Pods selected by the label,
admin@Jamess-MBP src %

```



```

admin@Jamess-MBP src % kubectl describe nodes | grep -i address -A 1
Addresses:
  InternalIP: 192.168.64.2
admin@Jamess-MBP src %

```

Internal ip of the nodes inside the cluster

private IP

```

admin@Jamess-MBP src % curl 192.168.64.2:32337

```

↳ lab VM is on the same machine as the node.
so will allow traffic.

will connect to webserver service.

we can try from any node & it will connect to that service.

(services expose pods via a static IP address.

(may be changing but service will take care of it

CONCLUSION

- Services expose Pods via a static IP address
- NodePort Services allows access from outside the cluster

Multi-container pods

Our Application

- A counter
- 4 containers split across 3 tiers
- Application tier is a Node.js server container
- Redis data tier storing the counter
- Poller and counter in the support tier
- All containers configured using environment variables



Kubernetes Namespaces

Namespaces separate resources according to users, environments, or applications

Role-based access control (RBAC) to secure access per Namespace

Using Namespaces is a best practice

```
admin@Jamess-MBP src % cat 3.1-namespace.yaml
apiVersion: v1
kind: Namespace
metadata:
  name: microservice
  labels:
    app: counter
admin@Jamess-MBP src % kubectl create -f 3.1-namespace.yaml
namespace/microservice created

admin@Jamess-MBP src % cat 3.2-multi_container.yaml
apiVersion: v1
kind: Pod
metadata:
  name: app
spec:
  containers:
    - name: redis
      image: redis:latest
      imagePullPolicy: IfNotPresent
      ports:
        - containerPort: 6379
    - name: server
      image: lrakai/microservices:server-v1
      ports:
        - containerPort: 8080
      env:
        - name: REDIS_URL
          value: redis://localhost:6379
    - name: counter
      image: lrakai/microservices:counter-v1
      env:
        - name: API_URL
          value: http://localhost:8080
    - name: poller
      image: lrakai/microservices:poller-v1
      env:
        - name: API_URL
          value: http://localhost:8080
```

```
admin@Jamess-MBP src % kubectl create -f 3.2-multi_container.yaml -n microservice
pod/app created

admin@Jamess-MBP src % kubectl get -n microservice pod app
NAME READY STATUS RESTARTS AGE
app 0/4 ContainerCreating 0 50s
admin@Jamess-MBP src %
```

→ namespace name.
→ always mention specific pod than label
→ better always mention specific label
→ we are not mentioning namespace in here because if we do that pod will become less portable.
environment variable
→ is each container in pod share same network stack.

→ mentioning name space.

describe gives more details than get
* once running we can see (stderr, stdout) logs in the container.

```
admin@Jamess-MBP src % kubectl logs -n microservice app counter --tail 10
Incrementing counter by 8 ...
Incrementing counter by 9 ...
Incrementing counter by 3 ...
Incrementing counter by 2 ...
Incrementing counter by 2 ...
Incrementing counter by 9 ...
Incrementing counter by 2 ...
Incrementing counter by 10 ...
Incrementing counter by 6 ...
Incrementing counter by 9 ...

admin@Jamess-MBP src % kubectl logs -n microservice app poller -f → to follow logs.
```

Issue with above approach.

- ↳ Kubernetes → pods are building blocks
- ↳ It can scale only by increasing number of pods not by increasing containers inside the pods.



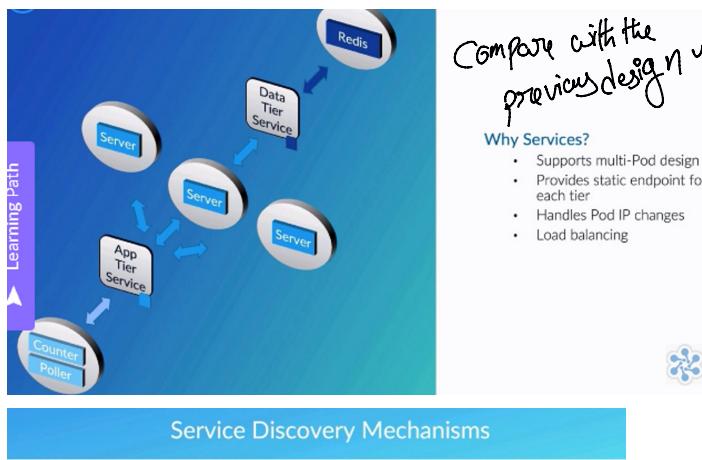
This is not what we want

They are tightly coupled!



↳ we should break application to multiple pods & then scale will be ideal.

Service discovery



Environment Variables

- Services address automatically injected in containers
- Environment variables follow naming conventions based on service name

DNS

- DNS records automatically created in cluster's DNS
- Containers automatically configured to query cluster DNS

```
admin@James-MBP src % kubectl create -f 4.1-namespace.yaml
namespace/service-discovery created
admin@James-MBP src % cat
```

```
admin@James-MBP src % cat 4.2-data_tier.yaml
apiVersion: v1
kind: Service
metadata:
  name: data-tier
  labels:
    app: microservices
spec:
  ports:
    - port: 6379
      protocol: TCP # default
      name: redis # optional when only 1 port
  selector:
    tier: data
    type: ClusterIP # default
  apiVersion: v1
  kind: Pod
  metadata:
    name: data-tier
    labels:
      app: microservices
      tier: data
  containers:
    - name: redis
```

Learning Path

type: cluster IP → default type, creates virtual ip inside cluster for internal access only.

```
in@James-MBP src % kubectl create -f 4.2-data_tier.yaml -n service-discovery
```

↳ resources are created in the order they are listed in the file.

↳ resources are created in the order they are listed in the file.

```
admin@Jamess-MBP src % kubectl describe service -n service-discovery data-tier
Name:           data-tier
Namespace:      service-discovery
Labels:         app=microservices
Annotations:   <none>
Selector:      tier=data
Type:          ClusterIP
IP:            10.98.110.202
Port:          redis 6379/TCP
TargetPort:    6379/TCP
Endpoints:     172.17.0.3:6379
Session Affinity: None
Events:        <none>
admin@Jamess-MBP src %
```

```
tier: app
apiVersion: v1
kind: Pod
metadata:
  name: app-tier
  labels:
    app: microservices
    tier: app
spec:
  containers:
    - name: server
      image: lrakai/microservices:server-v1
      ports:
        - containerPort: 8080
      env:
        - name: REDIS_URL
          # Environment variable service discovery
          # Naming pattern:
          # IP address: <all_caps_service_name>_SERVICE_HOST
          # Port: <all_caps_service_name>_SERVICE_PORT
          # Named Port: <all_caps_service_name>_SERVICE_PORT_<all_caps_port_name>
          value: redis://$(DATA_TIER_SERVICE_HOST):$(DATA_TIER_SERVICE_PORT_REDIS)
          # In multi-container example value was
          # value: redis://localhost:6379
in@Jamess-MBP src %
```

environment variables are auto created by the K8's
we just need naming convention proper.

- ex:- NAME_SERVICE_HOST

↳ data-tier

↳ hyphen replaced by underscore

ALL CAPS.

\$(DATA-TIER-SERVICE-HOST)

\$(DATA-TIER-SERVICE-PORT-REDIS)

Port name.

To use environment variables the service should be created first before
using/defining.

service should also be in same namespace where the env variables are used

```
admin@Jamess-MBP src % cat 4.4-support-tier.yaml
apiVersion: v1
kind: Pod
metadata:
  name: support-tier
  labels:
    app: microservices
    tier: support
spec:
  containers:
    - name: counter
      image: lrakai/microservices:counter-v1
      env:
        - name: API_URL
          # DNS for service discovery
          # Naming pattern:
          # IP address: <service_name>.<service_namespace>
          # Port: needs to be extracted from SRV DNS record
          value: http://$(DATA_TIER_SERVICE_HOST):$(DATA_TIER_SERVICE_PORT_API)
```

Support tier.

? let's create DNS for every service.

```

env:
  - name: API_URL
    # DNS for service discovery
    # Naming pattern:
    # IP address: <service_name>.<service_namespace>
    # Port: needs to be extracted from SRV DNS record
    value: http://app-tier.service-discovery:8080
  - name: poller
    image: lrakai/microservices:poller-v1
    env:
      - name: API_URL
        # omit namespace to only search in the same namespace
  - name: poller
    image: lrakai/microservices:poller-v1
    env:
      - name: API_URL
        # omit namespace to only search in the same namespace
        value: http://app-tier:${APP_TIER_SERVICE_PORT}

```

if in some namespace no need to mention.

continuation (or) use env. variable.
hard code (or) use env. variable.

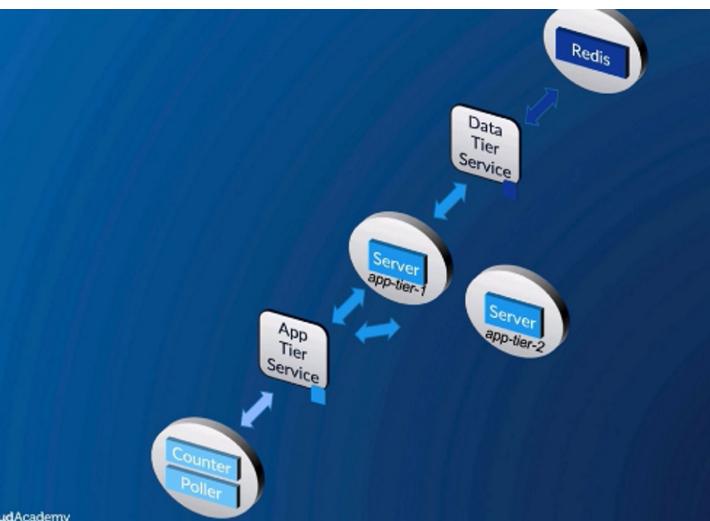
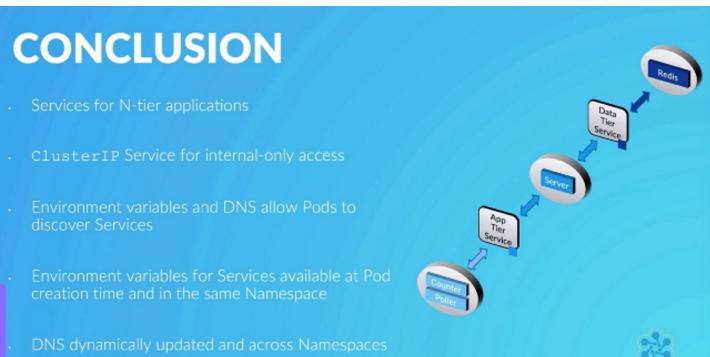
namespace omitted

DNS does not require underscore or all caps.

```

admin@Jamess-MBP src % kubectl create -f 4.4-support_tier.yaml -n service-discovery
pod/support-tier created
admin@Jamess-MBP src % kubectl get pods -n service-discovery
NAME          READY   STATUS    RESTARTS   AGE
-tier         1/1     Running   0          4m3s
-a-tier       1/1     Running   0          12m
port-tier    2/2     Running   0          11s

```



let me .

→ Till now we created pods directly. normally/ideally we should not create pod directly but instead use the higher level to create pods for example deployment.

Deployments

 Represent multiple replicas of a Pod

 Describe a desired state that Kubernetes needs to achieve

 Deployment Controller master component converges actual state to the desired state

We'll replace the individual pods with deployments that manage the pods for us.

Deployment is a template to create pods

↳ used to create replicas
↳ are copies of pods & scaled.

Comparing pod & deployment

```
protocol: TCP # default
  name: redis # optional when only 1 port
spec:
  selector:
    tier: data
  type: ClusterIP # default
---
apiVersion: v1
kind: Pod
metadata:
  name: data-tier
  labels:
    app: microservices
    tier: data
  annotations:
    - name: redis
      image: redis:latest
      imagePullPolicy: IfNotPresent
  ports:
    - containerPort: 6379
```

↳ no of pods:
↳ To group pods during deploy many.

→ instead of v1, it's apps/v1
and is deployment

→ spec → contains deployment specific settings
and also a pod template with exactly same as pod

does not need a name for pod here because
to create unique name for each pod.

```
admin@Jamess-MBP src % kubectl create -n deployments -f 2-data_tier.yaml -f 5.3-ap_p-tier.yaml -f 5.4-support_tier.yaml
```

↳ using multiple files at a time.

```
admin@Jamess-MBP src % kubectl get -n deployments deployments
NAME        READY   UP-TO-DATE   AVAILABLE   AGE
app-tier   1/1     1           1           76s
data-tier   1/1     1           1           76s
support-tier 1/1     1           1           76s
admin@Jamess-MBP src %
```

```
admin@Jamess-MBP src % kubectl -n deployments get pods
NAME          READY   STATUS    RESTARTS   AGE
app-tier-74cd4c68c9-z99h4  1/1     Running   0          2m24s
data-tier-8646dd765b-k6h98  1/1     Running   0          2m24s
support-tier-997bc57fb-c579t 2/2     Running   0          2m24s
```

↳ counter & poller
we can use scale command to scale the replicated.

```
admin@Jamess-MBP src % kubectl scale -n deployments deployments support-tier --replicas=5
deployment.apps/support-tier scaled
admin@Jamess-MBP src %
```

```
admin@Jamess-MBP src % kubectl -n deployments get pods
NAME          READY   STATUS    RESTARTS   AGE
app-tier-74cd4c68c9-z99h4  1/1     Running   0          4m49s
data-tier-8646dd765b-k6h98  1/1     Running   0          4m49s
support-tier-997bc57fb-42s98 2/2     Running   0          54s
```

↳ matching poller & counter remain same for pod

```

admin@Jamess-MBP ~ % kubectl -n deployments get pods
NAME          READY   STATUS    RESTARTS   AGE
app-tier-74cd4c68c9-h7hvf  1/1     Running   0          4m49s
data-tier-8646dd765b-k6h98  1/1     Running   0          4m49s
support-tier-997bc57fb-42s98 2/2     Running   0          54s
support-tier-997bc57fb-c579t 2/2     Running   0          4m49s
support-tier-997bc57fb-191kc 2/2     Running   0          54s
support-tier-997bc57fb-lp56v 2/2     Running   0          54s
support-tier-997bc57fb-mtfrc 2/2     Running   0          54s

```

replicated deployment pods, not individual containers inside the pod.

notice containers poller & counter remain same per pod because

Deployment ensure that specified number of replicated pods keep running

C) Test this by deleting pods.

```

admin@Jamess-MBP ~ % kubectl delete -n deployments pods support-tier-997bc57fb-mtfrc
(1) support-tier-997bc57fb-lp56v support-tier-997bc57fb-191kc
pod "support-tier-997bc57fb-mtfrc" deleted
pod "support-tier-997bc57fb-lp56v" deleted
pod "support-tier-997bc57fb-191kc" deleted

```

```
admin@Jamess-MBP ~ % watch -n 1 kubectl -n deployments get pods
```

Every 1.0s: kubectl -n deployments get pods

NAME	READY	STATUS	RESTARTS	AGE
app-tier-74cd4c68c9-h7hvf	1/1	Running	0	9m37s
data-tier-8646dd765b-wjnjid	1/1	Running	0	9m37s
support-tier-997bc57fb-7597k	2/2	Running	0	3m8s
support-tier-997bc57fb-79gsl	2/2	Running	0	7m59s
support-tier-997bc57fb-8wjnf	2/2	Terminating	0	7m59s
support-tier-997bc57fb-b2f9f	2/2	Terminating	0	3m8s
support-tier-997bc57fb-b8c4q	2/2	Running	0	19s
support-tier-997bc57fb-fggcx	2/2	Running	0	19s
support-tier-997bc57fb-hprh6	2/2	Terminating	0	3m8s
support-tier-997bc57fb-k8kfd	2/2	Running	0	19s

Every 1.0s: kubectl -n deployments get pods

NAME	READY	STATUS	RESTARTS	AGE
app-tier-74cd4c68c9-h7hvf	1/1	Running	0	10m
data-tier-8646dd765b-wjnjid	1/1	Running	0	10m
support-tier-997bc57fb-7597k	2/2	Running	0	3m36s
support-tier-997bc57fb-79gsl	2/2	Running	0	8m27s
support-tier-997bc57fb-d8c4q	2/2	Running	0	47s
support-tier-997bc57fb-fggcx	2/2	Running	0	47s
support-tier-997bc57fb-k8kfd	2/2	Running	0	47s

```
admin@Jamess-MBP ~ % kubectl scale -n deployments deployment app-tier --replicas=5
```

```

deployment.apps/app-tier scaled
deployment.apps/app-tier scaled
admin@Jamess-MBP ~ % kubectl -n deployments get pods
NAME          READY   STATUS    RESTARTS   AGE
app-tier-74cd4c68c9-kd7z4  1/1     Running   0          13s
app-tier-74cd4c68c9-qnrhh  1/1     Running   0          13s
app-tier-74cd4c68c9-rkbkh  1/1     Running   0          13s
app-tier-74cd4c68c9-thh25  1/1     Running   0          13s
app-tier-74cd4c68c9-z99hs  1/1     Running   0          9m48s
data-tier-8646dd765b-k6h98  1/1     Running   0          9m48s
support-tier-997bc57fb-2lptf 2/2     Running   0          3m4s
support-tier-997bc57fb-42s98 2/2     Running   0          5m53s
support-tier-997bc57fb-c579t 2/2     Running   0          9m48s
support-tier-997bc57fb-gpmj5 2/2     Running   0          3m3s
support-tier-997bc57fb-zfcrn 2/2     Running   0          3m3s
admin@Jamess-MBP ~ %

```

```

admin@Jamess-MBP ~ % kubectl describe -n deployments service app-tier
Name:           app-tier
Namespace:      deployments
Labels:         app=microservices
Annotations:   <none>
Selector:      tier=app
Type:          ClusterIP
IP:            10.11.185.244
Port:          <unset>  8080/TCP
TargetPort:    8080/TCP
Endpoints:    172.17.0.10:8080,172.17.0.11:8080,172.17.0.12:8080 + 2 more...
Session Affinity: None
Events:        <none>

```

all 5 replicas are tracked & load balanced in the service.

CONCLUSION

- Deployments to manage Pods in each tier
- Kubernetes ensures actual state matches desired state
- kubectl scale to scale number of replicas
- Services seamlessly support scaling
- Scaling is best with stateless Pods

→ even there is any failure or manual deletion it ensure to bring back to desired state

→ we can't scale the data tier in above example.

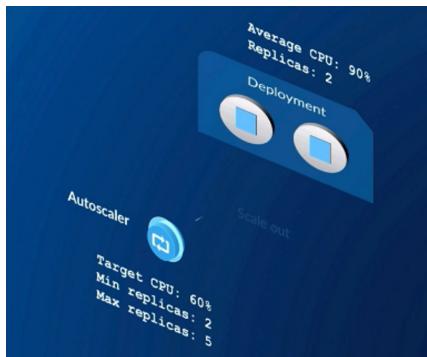
in

Autoscaling Deployments

 Scale automatically based on CPU utilization (or custom metrics)

 Set target CPU along with min and max replicas

 Target CPU is expressed as a percentage of the Pod's CPU request



Creates more replicas if average usage exceeds target.



Autoscaling depends on metrics being collected on the cluster.

Metrics

 Autoscaling depends on metrics being collected

 Metrics Server is one solution for collecting metrics

{ → we have to get metrics server up & running before using autoscaling.
Metrics server needed.



Metrics Server is one solution for collecting metrics



Several manifest files are used to deploy Metrics Server
(<https://github.com/kubernetes-sigs/metrics-server>)

```
admin@Jamess-MBP src % kubectl apply -f metrics-server/
clusterrole.rbac.authorization.k8s.io/system:aggregated-metrics-reader created
clusterrolebinding.rbac.authorization.k8s.io/metrics-server:system:auth-delegator created
rolebinding.rbac.authorization.k8s.io/metrics-server-auth-reader created
apiservice.apiregistration.k8s.io/v1beta1.metrics.k8s.io created
serviceaccount/metrics-server created
deployment.apps/metrics-server created
service/metrics-server created
clusterrole.rbac.authorization.k8s.io/system:metrics-server created
clusterrolebinding.rbac.authorization.k8s.io/system:metrics-server created
admin@Jamess-MBP src %
```

```
admin@Jamess-MBP src % kubectl top pods -n deployments
NAME          CPU(cores)   MEMORY(bytes)
app-tier-74cd4c68c9-h74qz  2m        45Mi
app-tier-74cd4c68c9-hcg26  1m        45Mi
app-tier-74cd4c68c9-p24sr  2m        45Mi
app-tier-74cd4c68c9-pstcr  2m        45Mi
app-tier-74cd4c68c9-sfxgm  2m        46Mi
data-tier-8646dd765b-7f457  2m        2Mi
support-tier-997bc57fb-2dtq2 10m       2Mi
support-tier-997bc57fb-dvlcv 11m       2Mi
support-tier-997bc57fb-gtd2d 12m       2Mi
support-tier-997bc57fb-kvvtr 11m       2Mi
support-tier-997bc57fb-mdlrb 12m       2Mi
admin@Jamess-MBP src %
admin@Jamess-MBP src %
```

```
--- previous
apiVersion: apps/v1
kind: Deployment
metadata:
  name: app-tier
  labels:
    app: microservices
    tier: app
spec:
  replicas: 1
  selector:
    matchLabels:
      tier: app
  template:
    metadata:
      labels:
        app: microservices
        tier: app
    spec:
      containers:
        - name: server
          image: lrakai/microservices:server-v1
          ports:
            - containerPort: 8080
      env:
        - name: REDIS_URL
          # Environment variable service discovery
          # Naming pattern:
          # IP address: <all_caps_service_name>.SERVICE_H
          # Port: <all_caps_service_name>.SERVICE_PORT
          # Named Port: <all_caps_service_name>.SERVICE_P
          value: redis://${DATA_TIER_SERVICE_HOST}:${DATA_T}
          # In multi-container example value was
          # value: redis://localhost:6379
admin@Jamess-MBP src %
```

```
--- updated
apiVersion: apps/v1
kind: Deployment
metadata:
  name: app-tier
  labels:
    app: microservices
    tier: app
spec:
  replicas: 5
  selector:
    matchLabels:
      tier: app
  template:
    metadata:
      labels:
        app: microservices
        tier: app
    spec:
      containers:
        - name: server
          image: lrakai/microservices:server-v1
          ports:
            - containerPort: 8080
          resources:
            requests:
              cpu: 20m # 20 milliCPU / 0.02 CPU
      env:
        - name: REDIS_URL
          # Environment variable service discovery
          # Naming pattern:
          # IP address: <all_caps_service_name>.SERVICE_H
          # Port: <all_caps_service_name>.SERVICE_PORT
          # Named Port: <all_caps_service_name>.SERVICE_P
          value: redis://${DATA_TIER_SERVICE_HOST}:${DATA_T}
          # In multi-container example value was
          # value: redis://localhost:6379
admin@Jamess-MBP src %
```

```
admin@Jamess-MBP src % kubectl create -f 6.1-app_tier_cpu_request.yaml -n deployments
Error from server (AlreadyExists): error when creating "6.1-app_tier_cpu_request.yaml": services "app-tier" already exists
Error from server (AlreadyExists): error when creating "6.1-app_tier_cpu_request.yaml": deployments.apps "app-tier" already exists
admin@Jamess-MBP src %
admin@Jamess-MBP src %
admin@Jamess-MBP src %
```

→ we can delete & recreate

(or)

→ use apply command which update the deployment.

* (KA) administrator course will have idea on diff between apply & create

```
admin@Jamess-MBP src % kubectl get -n deployments deployments app-tier
NAME      READY   UP-TO-DATE   AVAILABLE   AGE
app-tier  5/5     5           5           13m
admin@Jamess-MBP src %
```

J "autoscale"

all fully under metrics server
diff are being created.

1000 Mi = 1 CPU

K8S will deploy pods on to
each node having min of
of 20 milliCPU

```

admin@Jamess-MBP src % cat 6.2-autoscale.yaml
apiVersion: autoscaling/v1
kind: HorizontalPodAutoscaler
metadata:
  name: app-tier
  labels:
    app: microservices
    tier: app
spec:
  maxReplicas: 5
  minReplicas: 1
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: app-tier
  targetCPUUtilizationPercentage: 70

```

→ kind for autoscaling
 → lower & upper limit
 → 70%
 } command for above manifest —
 # Equivalent to
 # kubectl autoscale deployment app-tier --max=5 --min=1 --cpu-percent=70

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
app-tier	1/1	1	1	19m
data-tier	1/1	1	1	19m
support-tier	5/5	5	5	19m

admin@Jamess-MBP src % kubectl api-resources → gives all resources with their short names.

customresourcedefinitions	crd,crds	apiextensions.k8s.io	false	CustomResourceDefinition
apiservices		apiregistration.k8s.io	false	APIService
controllerrevisions		apps	true	ControllerRevision
demonsets	ds	apps	true	DaemonSet
deployments	deploy	apps	true	Deployment
replicasetss	rs	apps	true	ReplicaSet
statefulsets	sts	apps	true	StatefulSet
tokenreviews		authentication.k8s.io	false	TokenReview
localsubjectaccesreviews		authorization.k8s.io	true	LocalSubjectAccessReview
selfsubjectaccesreviews		authorization.k8s.io	false	SelfSubjectAccessReview
selfsubjectrulesreviews		authorization.k8s.io	false	SelfSubjectRulesReview
subjectaccesreviews		authorization.k8s.io	false	SubjectAccessReview
horizontalpodautoscalers	hpa	autoscaling	true	HorizontalPodAutoscaler
cronjobs		batch	true	CronJob
jobs		batch	true	Job
certificatesigningrequests	CSR	certificates.k8s.io	false	CertificateSigningRequest
leases		coordination.k8s.io	true	Lease
endpointslices		discovery.k8s.io	true	EndpointSlice
events	ev	events.k8s.io	true	Event

```

admin@Jamess-MBP src % kubectl describe -n deployments hpa
Name: app-tier
Namespace: default
Labels: app=microservices
          tier=app
Annotations: <none>
CreationTimestamp: Tue, 04 Aug 2020 12:10:16 +1200
Reference: Deployment/app-tier
Metrics:
  resource cpu on pods (as a percentage of request): 40% (8m) / 70%
Min replicas: 1
Max replicas: 5
Deployment pods: 1 current / 1 desired
Conditions:
  Type Status Reason Message
  ---- ---- ---- -----
  AbleToScale True ReadyForNewScale recommended size matches current size
  ScalingActive True ValidMetricFound the HPA was able to successfully calculate a replica count from cpu resource utilization (percent of request)
  ScalingLimited False DesiredWithinRange the desired count is within the acceptable range
Events:
  Type Reason Age From Message
  Normal SuccessfulRescale 107s horizontal-pod-autoscaler New size: 2; reason: All metrics below target
  Normal SuccessfulRescale 77s horizontal-pod-autoscaler New size: 1; reason: All metrics below target
admin@Jamess-MBP src %

```

```

admin@Jamess-MBP src % kubectl get -n deployments hpa
NAME      REFERENCE   TARGETS   MINPODS   MAXPODS   REPLICAS   AGE
app-tier  Deployment/app-tier  40%/70%  1         5           1          7m47s
admin@Jamess-MBP src %

```

admin@Jamess-MBP src %
 admin@Jamess-MBP src % kubectl edit -n deployments hpa
 (we can update manifest file & use apply)
 (or) (use edit command)
 Open in vi & we can edit the config like replicas if we want.

```

f:apiVersion: {}
f:kind: {}
f:type: {}
f:targetCPUUtilizationPercentage: {}
manager: kubectl
operation: Update
time: "2018-04-10T00:18:06Z"
- apiVersion: autoscaling/v1
  fieldsType: FieldsV1
  fieldsV1:
    f:metadata:
      f:annotations:
        :: {}
      f:autoscaling.alpha.kubernetes.io/conditions: {}
      f:autoscaling.alpha.kubernetes.io/current-metrics: {}
    f:status:
      f:currentCPUUtilizationPercentage: {}
      f:currentReplicas: {}
      f:desiredReplicas: {}
      f:lastScaleTime: {}
    manager: kube-controller-manager
    operation: Update
    time: "2018-04-10T00:18:07Z"
  name: app-tier
  namespace: deployments
  resourceVersion: "4191"
  selfLink: /apis/autoscaling/v1/namespaces/deployments/horizontalpodautoscalers/app-tier
  uid: d811baba-f4ac-432d-b5cc-67d6ee8535e62
spec:
  maxReplicas: 5
  minReplicas: 1
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: app-tier
  targetCPUUtilizationPercentage: 70
status:
  ...

```

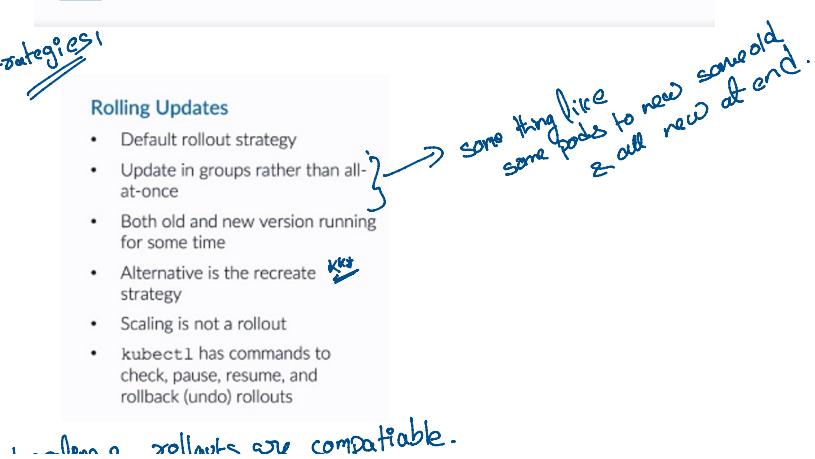
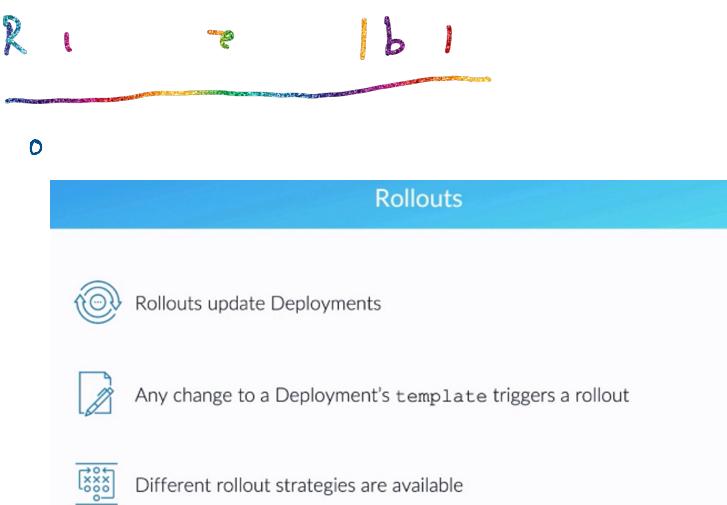
```

admin@James's-MBP ~ %
admin@James's-MBP ~ % kubectl edit -n deployments hpa
horizontalpodautoscaler.autoscaling/app-tier edited

```

Every 1.0s: kubectl get -n deployments deployment

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
app-tier	2/2	2	2	24m
data-tier	1/1	1	1	24m
support-tier	5/5	5	5	24m



check, pause, resume, and rollback (undo) rollouts

o autoscaling & rollouts are compatible.

```
Term2 Shell Edit View Session Scripts Profiles Toolbar Window Help
bash
admin@Jamess-MBP src % kubectl delete -n deployments hpa app-tier
horizontalpodautoscaler.autoscaling "app-tier" deleted
admin@Jamess-MBP src %
admin@Jamess-MBP src %
admin@Jamess-MBP src % kubectl edit -n deployments deployment app-ti...
```

```
Term2 Shell Edit View Session Scripts Profiles Toolbar Window Help
bash
{
  :message: {}
  :fireason: {}
  :status: {}
  :type: {}
  :forcePodEviction: {}
  :fireAndForget: {}
  :replicas: {}
  :updatedReplicas: {}
  manager: kube-controller-manager
  operation: Update
  time: "2020-08-04T03:01:31Z"
  namespaces: []
  resourceVersion: "1711"
  selflink: /apis/apps/v1/namespaces/deployments/deployments/app-tier
  uid: 0350bfcf-9bdf-4960-95a4-e43326b691a3
  spec:
    progressDeadlineSeconds: 600
    replicas: 10
    revisionHistoryLimit: 10
    selector:
      matchLabels:
        tier: app
      strategy:
        rollingUpdate:
          maxSurge: 25%
          maxUnavailable: 25%
        type: RollingUpdate
    status:
      desired:
        readyTimestamp: null
      available: 10
      app: microservices
      tier: app
    containers:
      env:
```

~ change replicas
from 2 to
10

```
resourceVersion: "1711"
selflink: /apis/apps/v1/namespaces/deployments/deployments/app-tier
uid: 0350bfcf-9bdf-4960-95a4-e43326b691a3
spec:
  progressDeadlineSeconds: 600
  replicas: 10
  revisionHistoryLimit: 10
  selector:
    matchLabels:
      tier: app
    strategy:
      rollingUpdate:
        maxSurge: 25%
        maxUnavailable: 25%
      type: RollingUpdate
  template:
    metadata:
      creationTimestamp: null
    labels:
      app: microservices
      tier: app
    containers:
      env:
        name: REDIS_URL
        value: redis://$(DATA_TIER_SERVICE_HOST):$(DATA_TIER_SERVICE_PORT_REDIS)
      image: lrokol/microservices:server-v1
      imagePullPolicy: IfNotPresent
      name: server
    ports:
      - containerPort: 8080
        protocol: TCP
    resources:
      requests:
        cpu: 20m
  terminationMessagePath: /dev/termination-log
  terminationMessagePolicy: File
```

3 delete the resources req. toward
any potential problems
while scheduling replicas.
If all is of the CPU,
requests can be satisfied.

```
admin@Jamess-MBP src % kubectl edit -n deployments deployment app-tier
deployment.apps/app-tier edited
admin@Jamess-MBP src %
admin@Jamess-MBP src %
admin@Jamess-MBP src % watch -n 1 kubectl get -n deployments deployment app-tier...
```

Every 1.0s: kubectl get -n deployments deployment app-tier

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
app-tier	10/10	10	10	8m25s

```
admin@Jamess-MBP src % kubectl edit -n deployments deployment app-tier
```

tier: app
strategy:
 rollingUpdate:
 maxSurge: 25%
 maxUnavailable: 25%
 type: RollingUpdate
template:

} (8) server added default values for the deployment strategy.
→ type is rollingupdate.
maxUnavailable → how many old pods can be deleted
... made to be ready.

type is rollingUpdate
 maxUnavailable → how many old pods can be deleted without waiting for new pods to be ready.
 maxSurge → how many replicas over the desired total are allowed during a rollout.

```

- name: REDIS_URL
  value: redis://${DATA_TIER_SERVICE_HOST}:${DATA_TIER_SERVICE_PORT_REDIS}
image: lrakai/microservices:server-v1
imagePullPolicy: IfNotPresent
name: cloucademy
ports:
- containerPort: 8080
  
```

→ changed name from server to cloud academy.
 → non-functional change.

```

admin@Jamesess-MBP src % kubectl edit -n deployments deployment app-tier
deployment.apps/app-tier edited
admin@Jamesess-MBP src %
admin@Jamesess-MBP src %
admin@Jamesess-MBP src % kubectl rollout -n deployments status deployment app-tier
Waiting for deployment "app-tier" rollout to finish: 6 out of 10 new replicas have been updated...
Waiting for deployment "app-tier" rollout to finish: 6 out of 10 new replicas have been updated...
Waiting for deployment "app-tier" rollout to finish: 7 out of 10 new replicas have been updated...
Waiting for deployment "app-tier" rollout to finish: 7 out of 10 new replicas have been updated...
Waiting for deployment "app-tier" rollout to finish: 7 out of 10 new replicas have been updated...
Waiting for deployment "app-tier" rollout to finish: 7 out of 10 new replicas have been updated...
Waiting for deployment "app-tier" rollout to finish: 7 out of 10 new replicas have been updated...
Waiting for deployment "app-tier" rollout to finish: 7 out of 10 new replicas have been updated...
Waiting for deployment "app-tier" rollout to finish: 8 out of 10 new replicas have been updated...
Waiting for deployment "app-tier" rollout to finish: 8 out of 10 new replicas have been updated...
Waiting for deployment "app-tier" rollout to finish: 8 out of 10 new replicas have been updated...
Waiting for deployment "app-tier" rollout to finish: 9 out of 10 new replicas have been updated...
Waiting for deployment "app-tier" rollout to finish: 9 out of 10 new replicas have been updated...
Waiting for deployment "app-tier" rollout to finish: 9 out of 10 new replicas have been updated...
Waiting for deployment "app-tier" rollout to finish: 9 out of 10 new replicas have been updated...
Waiting for deployment "app-tier" rollout to finish: 9 out of 10 new replicas have been updated...
Waiting for deployment "app-tier" rollout to finish: 9 out of 10 new replicas have been updated...
Waiting for deployment "app-tier" rollout to finish: 9 old replicas are pending termination...
Waiting for deployment "app-tier" rollout to finish: 3 old replicas are pending termination...
Waiting for deployment "app-tier" rollout to finish: 3 old replicas are pending termination...
Waiting for deployment "app-tier" rollout to finish: 2 old replicas are pending termination...
Waiting for deployment "app-tier" rollout to finish: 2 old replicas are pending termination...
Waiting for deployment "app-tier" rollout to finish: 1 old replicas are pending termination...
Waiting for deployment "app-tier" rollout to finish: 1 old replicas are pending termination...
Waiting for deployment "app-tier" rollout to finish: 1 old replicas are pending termination...
Waiting for deployment "app-tier" rollout to finish: 8 of 10 updated replicas are available...
Waiting for deployment "app-tier" rollout to finish: 9 of 10 updated replicas are available...
  
```

new replicas in }
old replicas out }

```

admin@Jamesess-MBP src % kubectl edit -n deployments deployment app-tier
deployment.apps/app-tier edited
admin@Jamesess-MBP src %
admin@Jamesess-MBP src %
admin@Jamesess-MBP src % kubectl rollout -n deployments pause deployment app-tier
deployment.apps/app-tier paused
admin@Jamesess-MBP src %
  
```

opened in two terminals to watch the deployment while we run pause command.

replicas which rollout before & pause continue to become ready }
Whereas the other which yet started will start after pause.

```

admin@Jamesess-MBP src % kubectl rollout -n deployments undo deployment app-tier
deployment.apps/app-tier rolled back
admin@Jamesess-MBP src %
  
```

to roll back the deployment.

```

admin@Jamesess-MBP src % kubectl rollout history -n deployments deployment app-tier
deployment.apps/app-tier
REVISION CHANGE-CAUSE
<none>
<none>
<none>
<none>
  
```

↳ to get all versions of deployment we can use & rollback to any version.

CONCLUSION

- Rollouts are triggered by Deployment template changes
- Rolling update is the default rollout strategy
- Rollouts can be paused, resumed, and undone

Kubernetes assumes that created containers are immediately ready and the rollout should continue. But this does not work in all cases. We may need to wait for the web server to accept connections. So here's another scenario. Considering an application using a relational database, the containers may start but it will fail until a database and tables are created.

These scenarios must be considered to build reliable applications. This is where probes in init containers come into the picture. So we'll take a look at integrating probes and init containers in our next two lessons.

b

- Rollout depend on containers. If rollout thinking that container is ready but there are cases where we have to wait like waiting for database.
- Another scenario is dead locked container & needed it to be restarted.

Probes are sometimes referred to health checks.

Readiness Probes



Used to check when a Pod is ready to serve traffic/handle requests



Useful after startup to check external dependencies



Readiness Probes set the Pod's ready condition. Services only send traffic to ready Pods.

to check if the pod is ready to handle requests

Ex:- to check connection to CMS is working otherwise we can't use token-adapter.

Liveness Probes



Detect when a Pod enters a broken state



Kubernetes will restart the Pod for you



Declared in the same way as readiness probes

Liveness Probes

- Readiness Probe → stop serving traffic.
- Liveness probe → restarts pod.

Declaring Probes

-  Probes can be declared in a Pod's containers
 -  All container probes must pass for the Pod to pass
 -  Probe actions can be a command that runs in the container, an HTTP GET request, or opening a TCP socket
 -  By default probes check containers every 10 seconds

We will add readiness and liveness probes to our application

- Liveness: Open TCP socket
 - Readiness: `redis-cli ping` command

App Tier (Server)

 - Liveness: HTTP GET /probe/liveness
 - Readiness: HTTP GET /probe/readiness

Learning Path

- Setting the initial delay seconds, we give the redis server an adequate time to start.
 - Given the consequences of failing a liveness probe is going to be restarting a Pod. It's generally advisable to have the liveness probe at a high delay than the readiness probe.
 - By default three sequential probes need to fail before a probe is marked as failed, so that we have some buffer. Kubernetes won't immediately restart the Pod the first time the probe fails, but we can configure it that way if we need to.
 - The particular delay depends on our application and how long it reasonably requires to start up. Five seconds should be more than enough to start checking readiness. And by default, we only need to pass a single probe before any traffic is sent to the Pod. Having the readiness initial delay too high will prevent Pods that are able to handle traffic from receiving any.

```
admin@James-MBP src % kubectl create -f 7.2-data-tier.yaml -n probes
service/data-tier created
deployment.apps/data-tier created
admin@James-MBP src %
admin@James-MBP src %
admin@James-MBP src % kubectl get deployments -n probes -w
NAME        READY   UP-TO-DATE   AVAILABLE   AGE
data-tier   0/1     1           0           2s
data-tier   1/1     1           1           14s
```

```

ubuntu@ip-10-0-128-5:~/Intro-to-k8s/src
$ kubectl get svc app-tier -o yaml
apiVersion: v1
kind: Service
metadata:
  name: app-tier
spec:
  selector:
    app: microservices
    tier: app
  ports:
    - port: 8080
      targetPort: 8080
      name: server
  clusterIP: None
  sessionAffinity: None
  externalTrafficPolicy: Cluster
status:
  loadBalancer: {}


  selector:
    app: microservices
    tier: app
  ports:
    - port: 8080
      targetPort: 8080
      name: server
  env:
    - name: REDIS_URL
      value: redis://localhost:6379

  # Environment variable service discovery
  # Naming pattern:
  # IP address: <all_cops_service_name>.SERVICE_H
  # Port: <all_cops_service_name>.SERVICE_PORT
  # Named Port: <all_cops_service_name>.SERVICE_P
  value: redis://$(DATA_TIER_SERVICE_HOST):$(DATA_T
  # In multi-container example value was
  # value: redis://localhost:6379

```

Annotations:

- dummy response { 200 } → checks with the server's readiness
- checks with the server's readiness { } → debug env variable optional
- makes sure correctly to database are working.

```

admin@James-MBP src % kubectl create -f 7.3-app-tier.yaml -n probes
service/app-tier created
deployment.apps/app-tier created
admin@James-MBP src %
admin@James-MBP src %
admin@James-MBP src % kubectl get deployments -n probes app-tier -w
NAME      READY   UP-TO-DATE   AVAILABLE   AGE
app-tier  1/1     1           1           25s

^Admin@James-MBP src %
admin@James-MBP src %
admin@James-MBP src % kubectl get -n probes pods
NAME        READY   STATUS    RESTARTS   AGE
app-tier-74c849f7c9-wdjn7  1/1     Running   0          65s
data-tier-7557d44554-l6g6h  1/1     Running   0          3m32s
admin@James-MBP src %
admin@James-MBP src %
admin@James-MBP src %

```

```

admin@James-MBP src %
admin@James-MBP src % kubectl logs -n probes app-tier-74c849f7c9-wdjn7 | cut -d' ' -f5,8-11_

```

```

03:33:52 urlencodedParser : /probe/liveness
03:33:56 dispatching GET /probe/readiness
03:33:56 query : /probe/readiness
03:33:56 expressInit : /probe/readiness
03:33:56 urlencodedParser : /probe/readiness
03:34:02 dispatching GET /probe/liveness
03:34:02 query : /probe/liveness
03:34:02 expressInit : /probe/liveness
03:34:02 urlencodedParser : /probe/liveness
03:34:06 dispatching GET /probe/readiness
03:34:06 query : /probe/readiness
03:34:06 expressInit : /probe/readiness
03:34:06 urlencodedParser : /probe/readiness
03:34:12 dispatching GET /probe/liveness
03:34:12 query : /probe/liveness
03:34:12 expressInit : /probe/liveness
03:34:12 urlencodedParser : /probe/liveness
03:34:16 dispatching GET /probe/readiness
03:34:16 query : /probe/readiness

```

CONCLUSION

- Container readiness probes monitor when Pods are ready to serve traffic and are temporarily out of service
- Container liveness probes monitor when Pods have entered a broken state and should be restarted
- Probes actions include commands, HTTP GET requests, and opening TCP sockets

n → If we want to test (it) prepare things before container starts we can do using init.

Motivation for Init Containers



Sometimes you need to wait for a service, downloads, dynamic or decisions before starting a Pod's containers



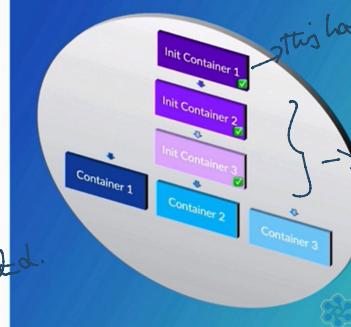
Prefer to separate initialization wait logic from the container image



Initialization is tightly coupled to the main application (belongs in the Pod)



Init containers allow you to run initialization tasks before starting the main container(s)



Init Containers

- Pods can declare any number of init containers
- Init containers run in order and to completion
- Use their own images
- Easy way to block or delay starting an application
- Run every time a Pod is created

Add an init container to our app tier that will wait for Redis before starting any application servers.

We'll see init containers have the same fields as regular containers except for `readinessProbe`.

```
laptop:src ~ % kubectl get pods -n probes
NAME          READY   STATUS    RESTARTS   AGE
app-tier-84db4f4665-rgt7f   1/1   Running   0          2m58s
admin@laptop:src ~ %
laptop:src ~ % kubectl describe pod app-tier-84db4f4665-rgt7f -n probes
existing init container
scripts which are already in the image & can be executed by the command here.

```

```
laptop:src ~ % kubectl get pods -n probes
NAME          READY   STATUS    RESTARTS   AGE
app-tier-84db4f4665-rgt7f   1/1   Running   0          2m58s
admin@laptop:src ~ %
laptop:src ~ % kubectl describe pod app-tier-84db4f4665-rgt7f -n probes

```

```
laptop:src ~ % kubectl get pods -n probes
NAME          READY   STATUS    RESTARTS   AGE
app-tier-84db4f4665-rgt7f   1/1   Running   0          2m58s
admin@laptop:src ~ %
laptop:src ~ % kubectl describe pod app-tier-84db4f4665-rgt7f -n probes
init container started.

```

```

$ cd /tmp/k8s
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...
$ kubectl logs -n probes app-tier-84db4f4665-rq37f -c default-redis
...

```

CONCLUSION

- Init containers allow you to perform tasks before main application containers have an opportunity to start
- Useful for checking preconditions and preparing dependencies

QUESTION

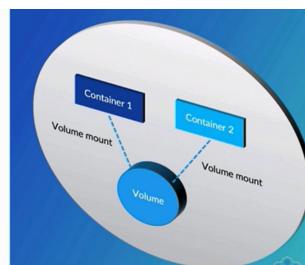
Motivation for Volumes

- Sometimes useful to share data between containers in a Pod } *in same Pod*
network is same
- Lifetime of container file systems is limited to the container's lifetime
- Can lead to unexpected consequences if a container restarts



Pod Storage in Kubernetes

- Two high-level storage options: Volumes and Persistent Volumes
- Used by mounting a directory in one or more containers in a Pod
- Pods can use multiple Volumes and Persistent Volumes
- Difference between Volumes and Persistent Volumes is how their lifetime is managed



Volumes

- Volumes are tied to a pod and their lifecycle
- Share data between containers and tolerate container restarts
- Use for non-persistent storage that is deleted with the Pod
- Default Volume type is emptyDir
- Data is lost if Pod is rescheduled on a different node

volume will be lost if pod is deleted (as node is changed)



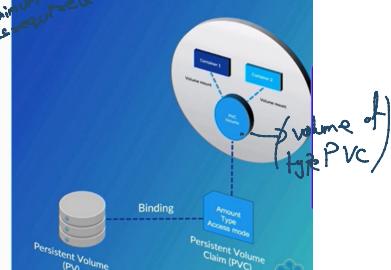
Persistent Volumes

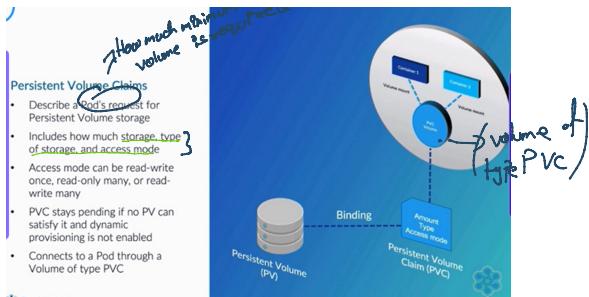
- Independent of Pod's lifetime
- Pods claim Persistent Volumes to use throughout their lifetime
- Can be mounted by multiple Pods on different Nodes if underlying storage supports it
- Can be provisioned statically in advance or dynamically on demand

How much provision size volume is required

Persistent Volume Claims

- Describe a Pod's request for Persistent Volume storage
- Includes how much storage_type of storage and access mode
- Access mode can be read-write once, read-only many, or read-write many
- PVC stays pending if no PV can satisfy it and dynamic provisioning is not enabled
- Connects to a Pod through a Volume of type PVC





Storage Volume Types

- Wide variety of volume types to choose from

 Use Persistent Volumes for more durable storage types

 Supported durable storage types include GCE Persistent Disks, Azure Disks, Amazon EBS, NFS, and iSCSI

Use a PersistentVolume for the sample applications data tier since we want the data to outlive its pod.

We'll use a statically provisioned Amazon Elastic Block Store (EBS) volume for the underlying storage.

```
NAME          READY   STATUS    RESTARTS   AGE
support-tier-748dd5cc5-fv1f   1/1     Running   0          42m
support-tier-599bc4fc8-8djh   1/1     Running   1          42m
support-tier-58d5d5458b-d8v8  2/2     Running   0          42m
kubernetes   1/1     Running   0          42m
```

clicking the namespace deployments
a load of support-tier.

```
NAME          READY   RESTARTS   AGE
app-pod-74d56cc54f   1/1     Running   1        44m
app-pod-599bf8c18dng   1/1     Running   1        44m
support-tier-58d5d5b6d  2/2     Running   0        44m

[check pods]

NAME          READY   RESTARTS   AGE
app-pod-74d56cc54f   1/1     Running   1        44m
app-pod-599bf8c18dng   1/1     Running   1        44m
support-tier-58d5d5b6d  2/2     Running   0        44m

[open bash of data layer]

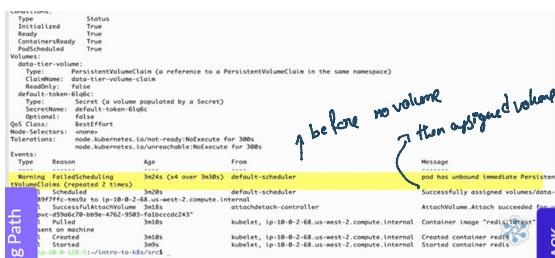
$ kubectl exec -n deployments data-tier-599bf8c18dng -it -- /bin/bash
root@data-tier-599bf8c18dng:/# kill 1
root@data-tier-599bf8c18dng:/# kill command terminated with exit code 137

[kill using process id]


```

```
ubuntu@ip-10-0-128-5:/intro-to-k8s/src$ cat 9.2-pv_data_tier.yaml | grep INSERT -CS  
capacity:
```

```
ubuntu@ip-10-0-128-5:~/intro-to-k8s/src$ cat 9.2-pv_data_tier.yaml | grep INSERT -C5
capacity:
  storage: 1Gi # 1 gibibyte
accessModes:
  - ReadWriteOnce
awsElasticBlockStore:
  volumeID: INSERT_VOLUME_ID # replace with actual ID
---
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: data-tier-volume-claim
ubuntu@ip-10-0-128-5:~/intro-to-k8s/src$ ubuntu@ip-10-0-128-5:~/intro-to-k8s/src$ sed -i "s/INSERT_VOLUME_ID/$vol_id/" 9.2-pv_data_tier.yaml
ubuntu@ip-10-0-128-5:~/intro-to-k8s/src$ ubuntu@ip-10-0-128-5:~/intro-to-k8s/src$ cat 9.2-pv_data_tier.yaml | grep $vol_id -C5
capacity:
  storage: 1Gi # 1 gibibyte
accessModes:
  - ReadWriteOnce
awsElasticBlockStore:
  volumeID: vol-0b071a7775995192c # replace with actual ID
---
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: data-tier-volume-claim
ubuntu@ip-10-0-128-5:~/intro-to-k8s/src$ 
```



```
root@ctrl-0-128-5:~/intro-to-k8s$ kubectl get pods -n volumes
NAME          READY   STATUS    RESTARTS   AGE
data-tier     1/1     Running   0          4m3s
support-tier  2/2     Running   0          4m3s
root@ctrl-0-128-5:~/intro-to-k8s$ kubectl logs -n volumes support-tier-68779db8-xvxt7
```

Every 5.0s: kubectl logs -n volumes support-tier-687789db8-xvxt7 poller --tail=1

```
Every 0.5s: kubectl logs -n volumes support-tier-687789db8-xvxt7 poller --tail=1  
Current counter: 825
```



CONCLUSION



Motivation for ConfigMaps and Secrets



Until now all container configuration has been in the Pod spec.



This makes it less portable than it could be



If sensitive information such as API keys and passwords is involved it presents a security issue



- specs
 - Results in easier to manage and more portable manifests
 - Both are similar but Secrets are specifically for sensitive data
 - There are specialized types of Secrets for storing Docker registry credentials and TLS certs
 - We will focus on the generic Secrets



Using ConfigMaps and Secrets



cd aimless-MBP- src % cat 10.4-app_tier_secret.yaml
apiVersion: v1
kind: Secret
type: Opaque

```
admin@Juniper-MBP:~/src$ k
admin@Juniper-MBP:~/src$ k
admin@Juniper-MBP:~/src$ kubectl exec -n config data-tier-b545b585-2lpcb -- redis-cli CONFIG GET tcp-keepalive
tcp-keepalive
0:0
admin@Juniper-MBP:~/src$ now it's changed
```

The screenshot shows a browser window with two tabs. The active tab is 'app-for-deployment' which contains a configuration file with annotations:

```
env:
  name: REDIS_URL
  # Environment variable service discovery
  value: redis://all-caps.service_name:SERVICE_PORT
    # IP:port = all-caps.service_name:SERVICE_PORT
  # Port: all-caps.service_name:SERVICE_PORT
  # Name: all-caps.service_name:SERVICE_NAME
  # Environment variable service discovery
  # value: redis://SCDATA.TIER_SERVICE_HOST:SCDATA_T
  # It means redis://localhost:6379
  # value: redis://localhost:6379
  - name: DB0
  - value: express*
```

Annotations in red highlight sections of the code:

- A box labeled "variable" highlights the first section of the code.
- A box labeled "variable" highlights the second section of the code.
- A box labeled "variable" highlights the third section of the code.
- A yellow box highlights the 'value' field of the 'REDIS_URL' variable.
- A yellow box highlights the 'value' field of the 'DB0' variable.
- A yellow box highlights the 'value' field of the 'REDIS_URL' variable.

The second tab is 'env' which lists environment variables:

Name	Value
REDIS_URL	# Environment variable service discovery
NamingConvention	# Naming Convention
PORT	# Port: all-caps.service_name:SERVICE_PORT
NAME	# Name: all-caps.service_name:SERVICE_NAME
ENVIRONMENT_VARIABLE	# Environment variable service discovery
SCDATA_T	value: redis://SCDATA.TIER_SERVICE_HOST:SCDATA_T
DB0	value: redis://localhost:6379

Annotations in red highlight specific values:

- A yellow box highlights the 'value' field of the 'REDIS_URL' variable.
- A yellow box highlights the 'value' field of the 'PORT' variable.
- A yellow box highlights the 'value' field of the 'NAME' variable.
- A yellow box highlights the 'value' field of the 'ENVIRONMENT_VARIABLE' variable.
- A yellow box highlights the 'value' field of the 'SCDATA_T' variable.
- A yellow box highlights the 'value' field of the 'DB0' variable.

A large yellow box highlights the 'value' field of the 'REDIS_URL' variable, with a note above it: "new variable (for configMap it's configMapKeyRef)" and a note below it: "name of the key, we get the value from".

```
admin@James-MBP:~/k8s$ kubectl create -f 10.5-app-tier.yaml -n config  
service/app-tier created  
deployment.apps/app-tier created  
admin@James-MBP:~/k8s$
```

} → created app tier.

- ConfigMaps and Secrets separate configuration data from pod specs or what would otherwise be stored in container images
- Both store data as key-value pairs
- Secrets are for sensitive data

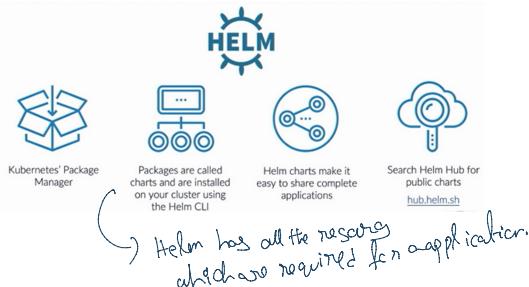
CONCLUSION

- ConfigMaps and Secrets separate configuration data from pod specs or what would otherwise be stored in container images
 - Both store data as key-value pairs
 - Secrets are for sensitive data

- Kubernetes Ecosystem**
- Vibrant ecosystem around the core of Kubernetes
 - We'll only touch on a select few members of the Kubernetes ecosystem



* CloudAcademy



Helm Examples

1. Use Helm for our application's Redis data tier
 - Redis charts available
 - Highly-available Redis chart avoids a single point of failure
 - Install with a single command
2. Create a chart for the entire sample application
 - Share it with anyone



Customize YAML manifests in Kubernetes



Helps you manage the complexity of your applications



Works by using a kustomization.yaml file that declares



Original manifests are untouched and remain usable



Using Kustomize



Kustomize is directly integrated with kubectl



Include the --kustomize or -k option to kubectl create or apply

Prometheus



Open-source monitoring and alerting system



A server for pulling in time series metric data and storing it



Inspired by an internal monitoring tool at Google called borgmon



De facto standard solution for monitoring Kubernetes

To some place where ICP came from



Prometheus + Kubernetes

- Kubernetes components supply all their own metrics in Prometheus format
 - Many more metrics than Metrics Server
- Adapter available to autoscale using metrics in Prometheus rather than CPU utilization
- Commonly paired with Grafana for visualizations (left image)
- Define alert rules and send notifications
- Easily installed via Helm chart



→ for machine learning
this is good.



Makes deployment of machine learning workflows on Kubernetes simple, scalable, and portable



A complete machine learning stack



Leverage Kubernetes to deploy anywhere, autoscale, etc.

Knative



Platform for building, deploying, and managing serverless workloads on Kubernetes



Can be deployed anywhere with Kubernetes, avoiding vendor lock-in



Supported by Google, IBM, and SAP

→ good for serverless.

I n



What Kubernetes Is

- Kubernetes is a distributed system
- Machines may be physical, virtual, on-prem, or in the cloud
- Schedules containers on machines
- Moves containers as machines are added/removed
- Can use different container runtimes
- Modular, extensible design



Kubernetes Architecture

- Cluster refers to all of the machines collectively and can be thought of as the entire running system
- Nodes are the machines in the cluster
- Nodes are categorized as workers or masters
- Worker nodes include software to run containers managed by the Kubernetes control plane
- Master nodes run the control plane
- The control plane is a set of APIs and software that Kubernetes users interact with
- The APIs and software are referred to as master components



Example `kubectl` commands

- `kubectl create` to create resources (Pods, Services, etc.)
- `kubectl delete` to delete resources
- `kubectl get` to get list of resources of a given type
 - `kubectl get pods`
- `kubectl describe` to print detailed info about a resource(s)
 - `kubectl describe pod server`
- `kubectl logs` to print container logs



Kubernetes in Practice

- Kubernetes terminology
- N-tier apps & Service discovery
- Deployment rollouts
- Monitoring containers with probes
- Preparing Pods with init containers
- Persistent data storage
- Separating configuration and sensitive data using ConfigMaps and Secrets





Kubernetes Ecosystem

- Vibrant ecosystem around the core of Kubernetes



Where To Go Next?

More Content on Cloud Academy

The image shows three learning path cards from Cloud Academy:

- Introduction to Kubernetes** by Logan Rakai. Tags: DevOps.
- Certified Kubernetes Administrator (CKA) Exam Preparation** by Logan Rakai. Tags: Deployment, DevOps.
- Certified Kubernetes Application Developer (CKAD) Exam Preparation** by Logan Rakai. Tags: Development, Deployment, DevOps.



Where To Go Next?

Convert your own application to run on Kubernetes



Take an application you know well and model it for Kubernetes



Reinforce what we learned and can help you find new solutions to existing challenges



Try packaging your solution as a Helm chart

Where To Go Next?

Participate in the Kubernetes Community



On GitHub
github.com/kubernetes



Join a Kubernetes Special Interest Group (SIG)



Conferences, Slack channels, and Google groups



KubeCon is the flagship Kubernetes conference